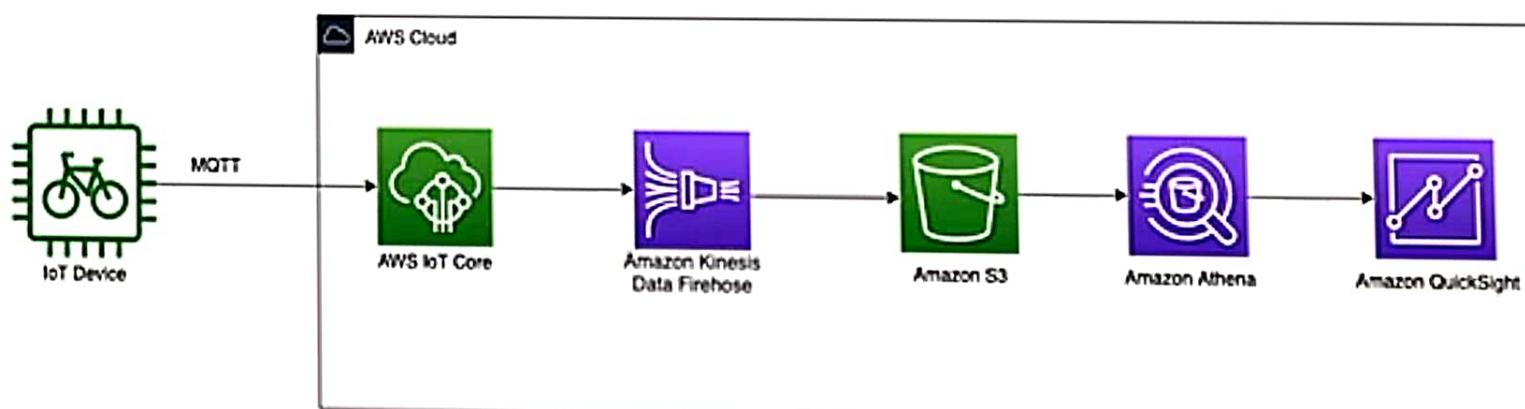


Pre-Requisites

- AWS Account
- 1 or more IoT devices sending messages to IoT Core

Pipeline to build



AWS Management Console

Stay connected to your AWS resources on-the-go

AWS Console Mobile App now supports four additional regions. Download the AWS Console Mobile App to your iOS or Android mobile device. [Learn More](#)

Explore AWS

AWS Certification

Get a complete overview of all things AWS Certification in this free e-book. [Learn more](#)

Free AWS Training

Complete projects faster and troubleshoot with confidence with 500+ free digital courses covering AWS products and services. [Learn more](#)

Get Started with Amazon S3

The most secure, durable, and scalable storage in the cloud. [Get started now](#)

AWS Certification

Discover the top 7 reasons to get AWS Certified. [Learn more](#)

Have feedback?

Submit Feedback to tell us about your experience with the AWS Management Console.

AWS services

Recently visited services

- IoT Core
- Amazon Translate
- Athena
- S3
- Lambda
- CloudFormation
- Amazon EventBridge
- VPC
- RDS
- API Gateway
- Support

All services

Build a solution

Get started with simple wizards and automated workflows.

| | | | |
|---|--|--|---|
| Launch a virtual machine With EC2 2-3 minutes | Build a web app With Elastic Beanstalk 6 minutes | Build using virtual servers With Lightsail 1-2 minutes | Register a domain With Route 53 3 minutes |
| Connect an IoT device With AWS IoT 5 minutes | Start migrating to AWS With AWS MGN 5-2 minutes | Start a development project With CodeStar 5 minutes | Deploy a serverless microservice With Lambda, API Gateway 2 minutes |

https://eu-west-1.console.aws.amazon.com/iot/home?region=eu-west-1&tab=home

AWS Services ▾

AWS IoT

Monitor

Activity

Onboard

Manage

Fleet Hub

Greengrass

Wireless connectivity

Secure

Defend

Art

Test

Device Advisor

MQTT test client

Software

Settings

Learn

Feature spotlight

Documentation

New console experience
Tell us what you think

Introducing the new AWS IoT console experience

We're updating the console experience for you. [Learn more](#) Try the new experiences and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT



AWS IoT is a managed cloud platform that lets connected devices - cars, light bulbs, sensor grids, and more - easily and securely interact with cloud applications and other devices.



Connect and manage your devices

Connect devices to the cloud using the protocol that best fits your requirements - HTTP, MQTT, or WebSocket. Devices can communicate with each other even if they are using different protocols.

[Learn more](#)



Process and act upon device data

Filter, transform, and act upon data from devices on the fly, based on business rules. AWS IoT can be easily integrated with AWS services like Amazon DynamoDB, Amazon Kinesis, Amazon Machine Learning, and AWS Lambda.

[Learn more](#)



Read and set device state at any time

AWS IoT stores the latest state of a device so that it can be read or set anytime, even when the device is offline.

[Learn more](#)

https://eu-west-1.console.aws.amazon.com/iot/home?region=eu-west-1&test

Search for services, features, marketplace products, and docs [Options+1]

AWS IoT Services ▾ Ireland Support ▾

AWS IoT

Monitor Activity Onboard Manage Fleet Hub Greengrass Wireless connectivity Secure DevLab Act Test Device Advisor MQTT test client

Introducing the new AWS IoT console experience. We're updating the console experience for you. Learn more Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > MQTT test client

MQTT test client

You can use the MQTT test client to monitor the MQTT messages being passed in your AWS account. Devices publish MQTT messages that are identified by topics to communicate their state to AWS IoT. AWS IoT also publishes MQTT messages to inform devices and apps of changes and events. You can subscribe to MQTT message topics and publish MQTT messages to topics by using the MQTT test client.

Subscribe to a topic Publish to a topic

Topic Filter info The Topic Filter describes the topicId to which you want to subscribe. The Topic Filter can include MQTT wildcards and characters.

edit tank/tester

Additional configuration

Subscribe

| Subscriptions | Topic |
|----------------------------------|--|
| You have no topic subscriptions. | Subscribe or select a topic to view incoming messages. |

New console experience Tell us what you think

Introducing the new AWS IoT console experience
We're updating the console experience for you. Learn more Try the new experiences and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > MQTT test client

MQTT test client

You can use the MQTT test client to monitor the MQTT messages being passed in your AWS account. Devices publish MQTT messages that are identified by topics to communicate their state to AWS IoT. AWS IoT also publishes MQTT messages to inform devices and apps of changes and events. You can subscribe to MQTT message topics and publish MQTT messages to topics by using the MQTT test client.

[Subscribe to a topic](#) [Publish to a topic](#)

Topic Filter: octank/escoter

The Topic Filter describes the topic(s) to which you want to subscribe. The Topic Filter can include MQTT wildcards characters.

octank/escoter

Additional configuration

[Subscribe](#)

| Subscriptions | octank/escoter | Pause | Clear | Export | Edit |
|----------------|--|-------|-------|--------|--------------------------------------|
| octank/escoter | <pre>[{"deviceID": "escoter28", "timestamp": 1629389252, "battery": 28, "speed": 70, "location": {"lat": 48.21438312333178, "long": 11.59591836431973}}]</pre> | | | | August 17, 2021, 11:54:12 (UTC+0200) |

New console experience Tell us what you think

Introducing the new AWS IoT console experience. We're updating the console experience for you. [Learn more](#). Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > MQTT test client

MQTT test client

You can use the MQTT test client to monitor the MQTT messages being passed in your AWS account. Devices publish MQTT messages that are identified by topics to communicate their state to AWS IoT. AWS IoT also publishes MQTT messages to inform devices and apps of changes and events. You can subscribe to MQTT message topics and publish MQTT messages to topics by using the MQTT test client.

[Subscribe to a topic](#) | [Publish to a topic](#)

Topic filter: `octank/escoter`

The topic filter describes the topic(s) to which you want to subscribe. The topic filter can include MQTT wildcards characters.

`octank/escoter`

Additional configuration

[Subscribe](#)

| Subscriptions | octank/escoter | Pause | Clear | Export | Edit |
|----------------|-----------------------------|-------|-------|--------|------|
| octank/escoter | <code>octank/escoter</code> | | | | |
| | <code>octank/escoter</code> | | | | |
| | { | | | | |
| | "deviceId": "escoter3", | | | | |
| | "timestamp": 1629389282, | | | | |
| | "battery": 42, | | | | |
| | "speed": 8, | | | | |
| | "location": { | | | | |
| | "lat": 48.13724994449474, | | | | |
| | "long": 11.58849885584588 | | | | |
| | } | | | | |
| | 1 | | | | |

August 17, 2021, 11:54:42 (UTC+0200)

New console experience
[Tell us what you think](#)

- Onboard
 - Manage
 - Fleet Hub
 - Greengrass
 - Wireless connectivity
 - Secure
 - Defend
- ▼ Act
- Rules**
 - Destinations
- ▼ Test
- Device Advisor
 - MQTT test client



You don't have any rules yet

Rules give your things the ability to interact with AWS and other web services. Rules are analyzed and actions are performed based on the messages sent by your things.

[Learn more](#)

[Create a rule](#)

Introducing the new AWS IoT console experience. We're updating the console experience for you. [Learn more](#). Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

Rule query statement
Indicate the source of the messages you want to process with this rule.
Using SQL version
2016-05-29

Rule query statement
SELECT <Attribute> FROM <Topic Filter> WHERE <Condition>. For example: SELECT temperature FROM 'iot/topic' WHERE temperature > 50. To learn more, see AWS IoT SQL Reference.

```
1 SELECT * FROM <Topic Filter>
```

Set one or more actions
Select one or more actions to happen when the above rule is matched by an inbound message. Actions define additional activities that occur when messages arrive, like storing them in a database, invoking cloud functions, or sending notifications. (* required)

Add action

Error action
Optionally set an action that will be executed when something goes wrong with processing your rule.

Add action

Tags

AWS IoT

Monitor

Activity

Onboard

Manage

Fleet Hub

Greengrass

Wireless connectivity

Secure

Device

Art

Roles

Destinations

Test

Device Advisor

MQTT test client

Software

Settings

Learn

Feature spotlight

Documentation

New console experience
Tell us what you think

Search for services, features, marketplace products, and docs [Option+I]

Introducing the new AWS IoT console experience
We're updating the console experience for you. Learn more Try the new experiences and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > Rules >

Select an action

Select an action:

- Insert a message into a DynamoDB table DYNAMODB
- Split message into multiple columns of a DynamoDB table (DynamoDBv2) DYNAMODB
- Send a message to a Lambda function LAMBDA
- Send a message as an SNS push notification SNS
- Send a message to an SQS queue SQS
- Send a message to an Apache Kafka cluster APACHE KAFKA
- Send a message to an Amazon Kinesis Stream AMAZON KINESIS
- Republish a message to an AWS IoT topic AWS IoT Broker

Monitor

Activity

► Onboard

► Manage

► Fleet Hub

► Greengrass

► Wireless connectivity

► Secure

► Devland

▼ IoT

► Roles

Destinations

▼ Test

► Device Advisor

MQTT test client

Software

Settings

Learn

Feature spotlight

Documentation

New console experience
Tell us what you think

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream
AMAZON-KINESIS-FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name:

Choose a resource

Separator (?)

No separator

Batch mode (?)

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Cancel

-  Send a message to the Amazon Elasticsearch Service
AMAZON ELASTICSEARCH
-  Send a message to a Salesforce IoT Input Stream
SALESFORCE IoT
-  Send a message to IoT Analytics
IOT ANALYTICS
-  Send a message to an IoT Events Input
IOT EVENTS
-  Send message data to asset properties in AWS IoT SiteWise
IOT SITESWISE
-  Start a Step Functions state machine execution
STEP FUNCTIONS
-  Send a message to a downstream HTTPS endpoint
HTTPS
-  Write a message into a Timeseries table
TIMESTREAM

Cancel

Configure action 

AWS IoT

Monitor

Activity

Onboard

Manage

Fleet Hub

Greengrass

Wireless connectivity

Secure

Defend

Act

Rules

Destinations

Test

Device Advisor

MQTT test client

Software

Settings

Learn

Feature spotlight

Documentation

New console experience
Tell us what you think

Introducing the new AWS IoT console experience
We're updating the console experience for you. Learn more. Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream

AMAZON KINESIS FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name

Choose a resource

Separator

No separator

Batch mode

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Cancel

The screenshot displays the AWS IoT 'Configure action' configuration interface. At the top, a blue header bar indicates the transition to the new console experience, with a 'Try the new experience' button and a 'Feedback' link. The main form is titled 'Configure action' and specifies the destination as 'Send a message to an Amazon Kinesis Firehose stream'. It includes fields for 'Stream name' (with a dropdown for 'Choose a resource' and a 'Create a new resource' button), 'Separator' (set to 'No separator'), and 'Batch mode' (unchecked). Below this, a section for selecting a role to grant AWS IoT access is shown, with a note that 'No role selected' and a 'Select' button. A 'Cancel' button and an 'Add action' button are at the bottom right. On the left, a sidebar lists various AWS IoT features like Monitor, Activity, and Rules, with 'Rules' currently selected. A 'New console experience' feedback link is also present in the sidebar.

Amazon Kinesis Data Firehose

Welcome to Amazon Kinesis Data Firehose

Amazon Kinesis Data Firehose is a fully managed, elastic service to easily deliver real-time data streams to destinations such as Amazon S3 and Amazon Redshift. You can start using Kinesis Data Firehose by:

1. Creating a delivery stream.
2. Sending your data to your delivery stream via Kinesis Agent or Kinesis Data Firehose APIs.

Data will be automatically delivered to your specified destination.

[Create Delivery Stream](#)

Kinesis Data Firehose Benefits



Easy to Use
Capture and deliver streaming data into destinations without writing any application or managing any infrastructure.



Direct to Data Stores
Batch, compress, and encrypt streaming data for delivery into your S3 bucket or Redshift cluster in as little as fifty seconds.



Zero Maintenance
Scale elastically to handle spikes in streaming data without intervention. Monitor the metrics for streaming data flowing into destinations.

Step 1: Name and source

- Step 2: Process records
- Step 3: Choose a destination
- Step 4: Configure settings
- Step 5: Review

New delivery stream

Delivery streams automatically and continuously load data to the specified destinations. Kinesis Data Firehose resources are not covered under the AWS Free Tier [?], and usage-based charges apply. For more information, see [Kinesis Data Firehose pricing](#) [?].

Delivery stream name

Acceptable characters are uppercase and lowercase letters, numbers, underscores, hyphens, and periods.

Choose a source

Choose how you would prefer to send records to the delivery stream.



Source

To learn about enabling server-side encryption (SSE), see [Data Protection in Amazon Kinesis Data Firehose](#) [?].

Direct PUT or other sources

Choose this option to send records directly to the delivery stream, or to send records from AWS IoT, CloudWatch Logs, or CloudWatch Events.

Kinesis Data Stream

Server-side encryption for source records in the delivery stream

Sales

Services ▾

Step 4: Configure settings

Step 5: Review

Delivery stream name: octane_deliverystream

Acceptable characters are uppercase and lowercase letters, numbers, underscores, hyphens, and periods.

Choose a source

Choose how you would prefer to send records to the delivery stream.

Firehose data flow overview

```
graph LR; Source[Source] --> Firehose[Firehose delivery stream]; Firehose --> Destination[Destination]
```

The diagram illustrates the data flow process. It starts with a 'Source' icon containing three document symbols. An arrow points from this to a central 'Firehose delivery stream' box. This box contains two sections: 'Source records' (with four document icons) and 'Processed records' (with five document icons). Another arrow points from the central box to a 'Destination' icon containing three document symbols. A dashed line labeled 'Optional' connects the 'Source' icon directly to the 'Destination' icon.

Source

To learn about enabling server-side encryption (SSE), see [Data Protection in Amazon Kinesis Data Firehose](#).

Direct PUT or other sources

Choose this option to send records directly to the delivery stream, or to send records from AWS IoT, CloudWatch Logs, or CloudWatch Events.

Kinesis Data Stream

Server-side encryption for source records in the delivery stream

You can use AWS Key Management Service (KMS) to create and manage customer keys (CMKs) and to control the use of encryption across a wide range of AWS services and in your applications.

Enable server-side encryption for source records in delivery stream

▶ [How to send source records to Kinesis Data Firehose](#)

Kinesis Data Firehose - Create delivery stream

Step 1: Name and source
Step 2: Process records
Step 3: Choose a destination
Step 4: Configure settings
Step 5: Preview

Choose a destination

[Learn more](#)

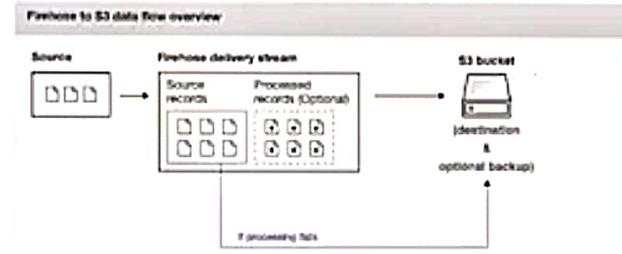
Destination

- Amazon S3**
Object storage built to store and retrieve any amount of data from anywhere.
- Amazon Redshift**
An enterprise-grade, petabyte-scale, fully managed data warehousing service.
- Amazon Elasticsearch**
An open-source search and analytics engine for use cases such as log analysis, real-time application monitoring, and wire-stream analysis.
- HTTP Endpoint**
A way to deliver data to your custom destinations.
- Third-party service provider**
Choose from a list of third-party service providers.

S3 destination

Choose a destination in Amazon S3 where your data will be stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.

[Learn more](#)



Q

amplify-undefined-dev-140225-deployment
amplify-undefined-dev-150120-deployment
amplify-undefined-dev-151124-deployment
aws-athena-query-results-eu-west-1-044118790602
aws-kendra-test-bucket
codeartifact-stage/bucket-4x2995090v
cf-templates-1a2bb061kcn-us-east-1
cf-templates-1a2bb061kcn-us-east-2
cloudtrail-awssage-044118790602-PrivateLog-Awsgend-dr-nets
delete
codepipeline-us-east-1-409130550901
config-bucket-044118790602
cloud-watch-gate-awsgarden-audit-044118790602
graviton2-pipeline-graviton2-pipeline/container-
hibernation
In-test-bucket
iot-device-simulator-iot-testdata-bucket-103101
iot-device-simulator-iot-testdata-bucket-263dry3n13w1
iot-device-simulator-iot-testdata-bucket-3111111111111111
Choose a bucket

S3 prefix

By default, Kinesis Data Firehose appends the prefix "YYYY/MM/dd/HH" (in UTC) to the data it delivers to Amazon S3. You can override this default by specifying a custom prefix that includes expressions that are evaluated at runtime.

If your custom prefix doesn't include expressions, Kinesis Data Firehose uses your prefix and appends "YYYY/MM/dd/HH". If your custom prefix includes a Kinesis Data Firehose random string or timestamp expression, Kinesis Data Firehose doesn't append "YYYY/MM/dd/HH". Learn more [\[?\]](#)

Backup S3 bucket prefix - optional

Enter a prefix

The screenshot shows the 'Choose a bucket' step in the AWS Kinesis Data Firehose configuration wizard. On the left, a dropdown menu displays a list of AWS buckets, including 'amplify-undefined-dev-140225-deployment', 'amplify-undefined-dev-150120-deployment', 'amplify-undefined-dev-151124-deployment', 'aws-athena-query-results-eu-west-1-044118790602', 'aws-kendra-test-bucket', 'codeartifact-stage/bucket-4x2995090v', 'cf-templates-1a2bb061kcn-us-east-1', 'cf-templates-1a2bb061kcn-us-east-2', 'cloudtrail-awssage-044118790602-PrivateLog-Awsgend-dr-nets', 'delete', 'codepipeline-us-east-1-409130550901', 'config-bucket-044118790602', 'cloud-watch-gate-awsgarden-audit-044118790602', 'graviton2-pipeline-graviton2-pipeline/container-hibernation', 'In-test-bucket', 'iot-device-simulator-iot-testdata-bucket-103101', 'iot-device-simulator-iot-testdata-bucket-263dry3n13w1', and 'iot-device-simulator-iot-testdata-bucket-3111111111111111'. Below the dropdown is a section titled 'S3 prefix' with explanatory text about the default prefix and how to override it with custom expressions. At the bottom, there is a field labeled 'Backup S3 bucket prefix - optional' with a placeholder 'Enter a prefix'.

S3 destination

Choose a destination in Amazon S3 where your data will be stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.

[Learn more](#)

Firehose to S3 data flow overview

```
graph LR; Source[Source] --> Firehose[Firehose delivery stream]; Firehose --> S3[S3 bucket]; S3 -- optional backup --> S3
```

The diagram illustrates the data flow from a source to a Firehose delivery stream, which then sends data to an S3 bucket. An optional backup path is shown from the S3 bucket back to itself.

S3 bucket

octank-database1

[Create new](#)

[View octank-database1 in S3 console](#)

S3 prefix

By default, Kinesis Data Firehose appends the prefix "YYYY/MM/dd/HH" (in UTC) to the data it delivers to Amazon S3. You can override this default by specifying a custom prefix that includes expressions that are evaluated at runtime.

If your custom prefix doesn't include expressions, Kinesis Data Firehose uses your prefix and appends "YYYY/MM/dd/HH". If your custom prefix includes a Kinesis Data Firehose random string or timestamp expression, Kinesis Data Firehose doesn't append "YYYY/MM/dd/HH". [Learn more](#)

Backup S3 bucket prefix - optional

Enter a prefix

S3 destination

Choose a destination in Amazon S3 where your data will be stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.

[Learn more](#)

Firehose to S3 data flow overview

```
graph LR; Source[Source] --> Firehose[Firehose delivery stream]; Firehose --> S3[S3 bucket]; S3 -- optional backup --> Firehose;
```

The diagram illustrates the data flow from a source to a Firehose delivery stream and finally to an S3 bucket. The Firehose delivery stream contains two sections: 'Source records' and 'Processed records (optional)'. An optional backup path is shown as a feedback loop from the S3 bucket back to the Firehose delivery stream.

S3 bucket

octank-database1

[View octank-database1 in S3 console](#)

S3 prefix

By default, Kinesis Data Firehose appends the prefix "YYYY/MM/dd/HH" (in UTC) to the data it delivers to Amazon S3. You can override this default by specifying a custom prefix that includes expressions that are evaluated at runtime.

If your custom prefix doesn't include expressions, Kinesis Data Firehose uses your prefix and appends "YYYY/MM/dd/HH". If your custom prefix includes a Kinesis Data Firehose random string or timestamp expression, Kinesis Data Firehose doesn't append "YYYY/MM/dd/HH". [Learn more](#)

Backup S3 bucket prefix - optional

Enter a prefix

S3 destination

Create a destination in Amazon S3 where your data will be stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.

Learn more ↗

Firehose to S3 data flow overview

```
graph LR; Source[Source] --> Firehose[Firehose delivery stream]; Firehose --> S3[S3 bucket];
```

The diagram illustrates the data flow from a source to a Firehose delivery stream, which then directs data to an S3 bucket. The Firehose delivery stream is shown with two parallel paths: "Source records" and "Processed records (optional)".

Create S3 bucket

A bucket is a container for objects stored in Amazon S3. Learn more

S3 bucket name
Region EU (Ireland)

S3 prefix

By default, Kinesis Data Firehose appends the prefix "YYYY/MM/DD/HH" (in UTC) to the data it delivers to Amazon S3. You can override this default by specifying a custom prefix that includes expressions that are evaluated at runtime.

If your custom prefix doesn't include expressions, Kinesis Data Firehose uses your prefix and appends "YYYY/MM/DD/HH". If your custom prefix includes a Kinesis Data Firehose random string or timestamp expression, Kinesis Data Firehose doesn't append "YYYY/MM/DD/HH". Learn more ↗

Backup S3 bucket prefix - optional
Enter a prefix

Services ▾

Search for services, features, marketplace products, and docs [Optimize ▾]

Amazon Kinesis Data Firehose | Documentation | Support

S3 destination

Choose a destination in Amazon S3 where your data will be stored. Amazon S3 is object storage built to store and retrieve any amount of data from anywhere.

[Learn more ▾](#)

Firehose to S3 data flow overview

```
graph LR; Source[Source] --> Firehose[Firehose delivery stream]; Firehose --> S3[S3 bucket]; S3 -- optional backup --> Firehose;
```

The diagram illustrates the data flow from a source to a Firehose delivery stream and then to an S3 bucket. The source feeds into the Firehose delivery stream, which contains 'Source records' and 'Processed records (Optional)'. The processed records can be used for 'If processing fails'. The output then goes to the S3 bucket, which is labeled as the 'destination' and includes an 'optional backup' feature.

S3 bucket

octank-datasink1

[View octank-datasink1 in S3 console](#) [?]

S3 prefix

By default, Kinesis Data Firehose appends the prefix "YYYY/MM/dd-HH" (in UTC) to the data it delivers to Amazon S3. You can override this default by specifying a custom prefix that includes expressions that are evaluated at runtime.

If your custom prefix doesn't include expressions, Kinesis Data Firehose uses your prefix and appends "YYYY/MM/dd-HH". If your custom prefix includes a Kinesis Data Firehose random string or timestamp expression, Kinesis Data Firehose doesn't append "YYYY/MM/dd-HH". [Learn more](#) [?]

Backup S3 bucket prefix - optional

Enter a prefix

Processing fails

S3 bucket
octank-database1

[View octank-database1 in S3 console](#) (?)

S3 prefix
By default, Kinesis Data Firehose appends the prefix "YYYY/MM/dd/HH" (in UTC) to the data it delivers to Amazon S3. You can override this default by specifying a custom prefix that includes expressions that are evaluated at runtime.
If your custom prefix doesn't include expressions, Kinesis Data Firehose uses your prefix and appends "YYYY/MM/dd/HH". If your custom prefix includes a Kinesis Data Firehose random string or timestamp expression, Kinesis Data Firehose doesn't append "YYYY/MM/dd/HH". Learn more (?)

Backup S3 bucket prefix - optional

S3 error prefix
You can specify an S3 bucket prefix to be used in error conditions. This prefix can include expressions for Kinesis Data Firehose to evaluate at runtime. Learn more about the rules for specifying prefix expressions (?)

Backup S3 bucket error prefix - optional



Kinesis Data Firehose - Create delivery stream

Step 1: Name and source

Step 2: Process records

Step 3: Choose a destination

Configure settings

Configure buffer, compression, logging, and IAM role settings for your delivery stream. [Learn more](#)

S3 buffer conditions

Kinesis Data Firehose buffers incoming records before delivering them to your S3 bucket. Record delivery will be triggered once either of these conditions has been satisfied. [Learn more](#)

Buffer size

MB

Enter a buffer size between 1 - 100 MB.

Buffer interval

seconds

Enter a buffer interval between 60 - 800 seconds.

S3 compression and encryption

Kinesis Data Firehose can compress records before delivering them to your S3 bucket. Compressed records can also be encrypted in the S3 bucket using a KMS master key. [Learn more](#)

S3 compression

- Disabled
- GZIP
- Snappy
- Zstd
- Hadoop-Compatible Snappy

S3 encryption

- Disabled
- Enabled

Salesforce Services ▾

Search for services, features, marketplace products, and docs [Option+I]

Enabled

Error logging

Kinesis Data Firehose data log record delivery errors to CloudWatch Logs. If enabled, a CloudWatch log group and corresponding log streams are created on your behalf. [Learn more](#)

Enabled

Tags - optional

You can add tags to organize your AWS resources, track items, and control access. [Learn more](#)

Key Value - optional

Enter key Enter value

You can add 49 more tag(s).

Permissions

IAM role

Kinesis Data Firehose uses this IAM role for all the permissions that the delivery stream needs. To specify different roles for the different permissions, use the API or the CLI. [Learn more](#)

Create or update IAM role **KinesisFirehoseServiceRole-octant_delivery-eu-west-1-1629180207571**

This creates the role or updates it if it already exists, sets the required policies to it, and enables Kinesis Data Firehose to assume it.

Choose existing IAM role

The role that you choose must have policies that include the permissions that Kinesis Data Firehose needs.

Cancel Previous



Review

Review your configuration details before creating your delivery stream.

Name and source

[Edit](#)

Delivery stream name

octank_deliverystream

Source

Direct PUT or other sources

Server-side encryption for source records

Disabled

Process records

[Edit](#)

Source record transformation

Disabled

Record format conversion

Disabled

Destination

[Edit](#)

Destination

Amazon S3

S3 bucket

octank-database1

S3 bucket Prefix

Destination

Destination
Amazon S3

S3 bucket
octane-database1-07

S3 bucket Prefix
-

S3 bucket error prefix
-

Settings

S3 buffer conditions
5 MB or 200 seconds

S3 compression
Disabled

S3 encryption
Disabled

Error logging
Enabled

Tags
No tags specified

IAM role
KinesisIngestServiceRole-octane-database-eu-west-1-1E2919308751

Amazon Kinesis

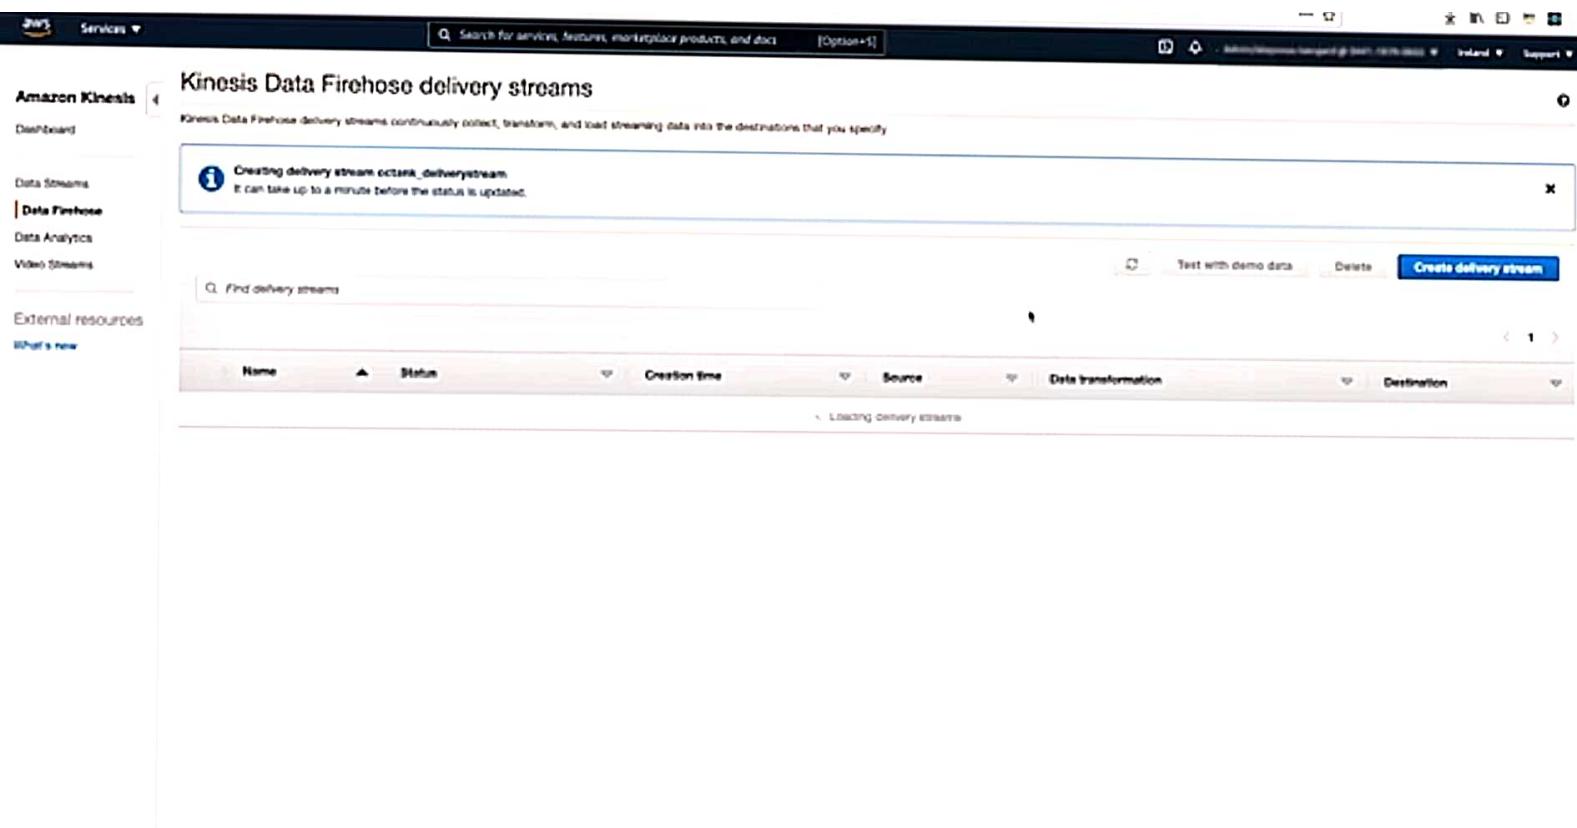
Kinesis Data Firehose delivery streams continuously collect, transform, and load streaming data into the destinations that you specify.

Creating delivery stream octank_deliverystream
It can take up to a minute before the status is updated.

Find delivery streams

Name Status Creation time Source Data transformation Destination

| Name | Status | Creation time | Source | Data transformation | Destination |
|-----------------------|----------|----------------------|------------------------------|---------------------|-----------------------------------|
| octank_deliverystream | Creating | 2021-08-17T11:39:02Z | Direct PUT and other sources | Disabled | Amazon S3 octank-database1 [7] |



Amazon Kinesis

Kinesis Data Firehose delivery streams continuously collect, transform, and load streaming data into the destinations that you specify.

Creating delivery stream octank_deliverystream
It can take up to a minute before the status is updated.

Find delivery streams

Name Status Creation time Source Data transformation Destination

| Name | Status | Creation time | Source | Data transformation | Destination |
|-----------------------|----------|----------------------|------------------------------|---------------------|-----------------------------------|
| octank_deliverystream | Creating | 2021-08-17T11:39:02Z | Direct PUT and other sources | Disabled | Amazon S3 octank-database1 [7] |

Sales Services ▾

Amazon Kinesis

Dashboard

Data Streams

Data Firehose

Data Analytics

Video Streams

External resources

What's new

Q Search for services, features, marketplace products, and docs [Option+I]

Kinesis Data Firehose delivery streams continuously collect, transform, and load streaming data into the destinations that you specify.

Successfully created delivery stream octank_deliverystream

Next, send records directly to the delivery stream using the Amazon Kinesis Agent or the Kinesis Data Firehose API using the AWS SDK, or send records from AWS IoT, CloudWatch Logs, or CloudWatch Events. Learn more

Q Find delivery streams

Test with demo data Delete Create delivery stream

| Name | Status | Creation time | Source | Data transformation | Destination |
|-----------------------|--------|----------------------|------------------------------|---------------------|-----------------------------------|
| octank_deliverystream | Active | 2021-08-17T11:29:02Z | Direct PUT and other sources | Disabled | Amazon S3 octank-database1 [7] |

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream AMAZON KINESIS FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name

Separator 

Batch mode

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Select

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream AMAZON-KINESIS-FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name

Separator 



Batch mode  

Choose or create a role to grant AWS IoT access to perform this action.



 Cancel

 Add action

AWS IoT

Monitor

Activity

Onboard

Manage

Fleet Hub

Greengrass

Wireless connectivity

Secure

Defend

Rules

Destinations

Test

Device Advisor

MQTT test client

Software

Settings

Learn

Feature spotlight

Documentation

New console experience
Tell us what you think

Search for services, features, marketplace products, and docs

Introducing the new AWS IoT console experience

We're updating the console experience for you. Learn more. Try the new experiences and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name:

octank_drivenstream

Create a new resource

Separator:

No separator

Batch mode:

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Create Role

Select

Cancel

Add action

AWS IoT

Introducing the new AWS IoT console experience
We're updating the console experience for you. Learn more Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > Rules >

Configure action

 Send a message to an Amazon Kinesis Firehose stream
AMAZON KINESIS FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name: octank_deliverystream

Separator: No separator

Batch mode:

Choose or create a role to grant AWS IoT access to perform this action.

All role selected

Cancel

New console experience Tell us what you think

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name

octank_deliverystream 

[Create a new resource](#)

Separator 

No separator 

Batch mode  

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

[Create Role](#)

[Select](#)

[Cancel](#)

[Add action](#)

AWS IoT

Introducing the new AWS IoT console experience. We're updating the console experience for AWS IoT.

Monitor Activity Onboard Manage Fleet Hub Greengrass Wireless connectivity Secure Defend Art Roles Destinations Test Device Advisor MQTT test client Software Settings Learn Feature spotlight Documentation [2]

New console experience Tell us what you think

Create a new role

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name: octank_deliverystream

Separator: No separator

Batch mode:

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Cancel

the navigation menu

AWS IoT

Introducing the new AWS IoT console experience
We're updating the console experience for you. Learn more [?] Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream

AMAZON KINESIS FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name: octank_deliverystream

Separator:

Batch mode:

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Cancel

New console experience
Tell us what you think [?]

Salesforce Home | Help & Support | Log In

AWS IoT Services Search for services, features, documentation, products, and more

AWS IoT Rules Create a new role

Configure action

Send a message to an Amazon Kinesis Firehose stream

Create role

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name: octank_deliverystream

Separator: No separator

Batch mode:

Choose or create a role to grant AWS IoT access to perform this action.

No role selected

Cancel

New console experience Tell us what you think

AWS IoT > Rules >

Configure action

Send a message to an Amazon Kinesis Firehose stream AMAZON KINESIS FIREHOSE

This will send the message to an Amazon Kinesis Firehose stream.

*Stream name

octank_deliverystream

Create a new resource

Separator

No separator

Batch mode

Choose or create a role to grant AWS IoT access to perform this action.

awsiotfirehose-role

Create Role

Select

Cancel

Add action

Set one or more actions

Select one or more actions to happen when the above rule is matched by an inbound message. Actions define additional activities that occur when messages arrive, like storing them in a database, invoking cloud functions, or sending notifications. (* required)



Send a message to an Amazon Kinesis Firehose stream
`event_delivery`

[Remove](#) [Edit](#)

[Add action](#)

Error action

Optionally set an action that will be executed when something goes wrong with processing your rule.

[Add action](#)

Tags

Apply tags to your resources to help organize and identify them. A tag consists of a case-sensitive key-value pair. [Learn more about tagging your AWS resources](#).

Tag name

Value

[Clear](#)

[Add another](#)

[Cancel](#)

[Create rule](#)

AWS Services ▾

AWS IoT

- Monitor
- Activity
- Onboard
- Manage
- Fleet Hub
- Greengrass
- Wireless connectivity
- Secure
- Defend
- Rules **Selected**
- Destinations
- Test
- Device Advisor
- MQTT test client

Software

Settings

Learn

Feature spotlight

Documentation

New console experience
Tell us what you think

Q Search for services, features, marketplace products, and docs [Option+I]

Introducing the new AWS IoT console experience
We're updating the console experience for you. Learn more Try the new experience and let us know what you think. You can turn off the new experience from the navigation menu.

Success
Successfully created rule.

AWS IoT > Rules

Rules

Create

| Name | Status |
|---------------------|--------|
| OctankAnalyticsRule | |

S https://eu-west-1.console.aws.amazon.com/iot/home?region=eu-west-1&rulehub

AWS Services AWS IoT

Search results for S3

Services (7) See all 7 results ▾

- S3** Scalable Storage in the Cloud
- S3 Glacier** Archive Storage in the Cloud
- Athena** Query Data in S3 using SQL
- AWS Snow Family** Large Scale Data Transport

Features (10) See all 10 results ▾

- Amazon S3 File Gateway** Storage Gateway feature
- Datasets** IoT Analytics feature
- Batch Operations** S3 feature
- Buckets** S3 feature

Documentation (401,673) See all 401,673 results in Documentation ▾

Create

New console experience Tell us what you think

This screenshot shows the AWS IoT console search results for "S3". The search bar at the top contains the query "S3". The left sidebar lists various AWS services and features under categories like Monitor, Activity, Onboard, Manage, Fleet Hub, Greengrass, Wireless connectivity, Secure, Behind, Art, Test, Device Advisor, MQTT test client, Software, Settings, Learn, Feature spotlight, and Documentation. The main content area displays three sections: Services, Features, and Documentation. The Services section lists S3, S3 Glacier, Athena, and AWS Snow Family. The Features section lists Amazon S3 File Gateway, Datasets (IoT Analytics feature), Batch Operations (S3 feature), and Buckets (S3 feature). The Documentation section links to over 400,000 results. A blue banner at the top right encourages users to "Tell us what you think" about the new console experience.

Amazon S3

We're continuing to improve the S3 console to make it faster and easier to use. If you have feedback on the updated experience, [Choose Provide Feedback](#).

Amazon S3

Buckets

Access Points

Object Lambda Access Points

Batch Operations

Access analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

AWS Organizations settings

Feature spotlight

AllS Marketplace for S3

Amazon S3

Account snapshot

Storage lens provides visibility into storage usage and activity trends. [Learn more](#)

View Storage Lens dashboard

Buckets

Buckets are containers for data stored in S3. [Learn more](#)

Search buckets by name

| Name | AWS Region | Access | Creation date |
|-----------------|------------|--------|---------------|
| Leading buckets | | | |

Amazon S3

We're continuing to improve the S3 console to make it faster and easier to use. If you have feedback on the updated experience, [Choose Provide Feedback](#).

Amazon S3

Amazon S3

Bucket

Access Points

Object Lambda Access Points

Batch Operations

Access analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

AWS Organizations settings

Feature spotlight

All AWS Marketplace for S3

Search for services, features, marketplace products, and docs [Options]

View Storage Lens dashboard

Buckets (24)

Buckets are containers for data stored in S3. Learn more.

Name AWS Region Access Creation date

| Name | AWS Region | Access | Creation date |
|--|---------------------------------|--------|---|
| amplify-undefined-dev-145121-deployment | US East (N. Virginia) us-east-1 | | December 31, 2020, 14:50:33 (UTC+01:00) |
| amplify-undefined-dev-150120-deployment | US East (N. Virginia) us-east-1 | | December 31, 2020, 15:10:28 (UTC+01:00) |
| amplify-undefined-dev-151014-deployment | US East (N. Virginia) us-east-1 | | December 31, 2020, 15:10:21 (UTC+01:00) |
| aws.athena-query-results-eu-west-1-044118790602 | EU (Ireland) eu-west-1 | | May 5, 2021, 18:24:37 (UTC+02:00) |
| ay-kendra-test-bucket | EU (Ireland) eu-west-1 | | June 28, 2021, 16:27:58 (UTC+02:00) |
| cdktest8-stagingbucket-4ad89909-5cav | EU (Ireland) eu-west-1 | | July 5, 2021, 12:42:11 (UTC+02:00) |
| cf-templates-1r2zbed94cm-us-east-1 | US East (N. Virginia) us-east-1 | | November 11, 2020, 11:52:10 (UTC+01:00) |
| cf-templates-1r2zbed94cm-us-east-2 | US East (Ohio) us-east-2 | | December 5, 2020, 14:22:19 (UTC+01:00) |
| cloudtrail-auditlog-044118790602-fvwsdpl-1engard-do-not-delete | US East (N. Virginia) us-east-1 | | October 20, 2020, 19:01:47 (UTC+02:00) |
| comipipeline-us-east-1-609130550931 | US East (N. Virginia) us-east-1 | | December 15, 2020, 14:47:55 (UTC+01:00) |
| config-bucket-044118790602 | US East (N. Virginia) us-east-1 | | January 25, 2021, 17:02:01 (UTC+01:00) |
| do-not-delete-patedgarden-audit-044118790602 | US West (Oregon) us-west-2 | | October 27, 2020, 01:20:06 (UTC+01:00) |
| gravitonid-pipeline-graviton2/pipelinelabcontainer-2zbogn4gpm | EU (Ireland) eu-west-1 | | July 5, 2021, 12:51:51 (UTC+02:00) |
| fr-test-bucket | EU (Ireland) eu-west-1 | | August 2, 2021, 17:46:58 (UTC+02:00) |
| iot-device-simulator-lambda-database-11klmuf5605am | US East (N. Virginia) us-east-1 | | February 2, 2021, 17:13:28 (UTC+01:00) |
| iot-device-simulator-lambda-database-26drq4ln15rl | EU (Ireland) eu-west-1 | | February 2, 2021, 17:21:15 (UTC+01:00) |
| iot-device-simulator-lambda-webdb-bucket-1fba0jdekr4 | US East (N. Virginia) us-east-1 | | February 2, 2021, 17:21:27 (UTC+01:00) |

Amazon S3

Amazon S3 > octank-datalake1

octank-datalake1

Objects Properties Permissions Metrics Management Access Points

Objects [1]

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Action Copy URL Copy URN Download Open Delete Actions Create folder Upload

Find objects by prefix:

| Name | Type | Last modified | Size | Storage class |
|-------|--------|---------------|------|---------------|
| 2021/ | Folder | - | - | - |

Actions

Provide Feedback

Block Public Access settings for this account

Storage Lens

Dashboards

AWS Organizations settings

Feature spotlight

AWS Marketplace for S3

The screenshot shows the AWS S3 console interface. On the left, there's a sidebar with various navigation links like 'Objects', 'Properties', 'Permissions', 'Metrics', 'Management', and 'Access Points'. The main area is titled 'octank-datalake1' and shows a single object named '2021/'. There are buttons for 'Upload' and 'Actions' (with options like 'Copy URL', 'Copy URN', 'Download', 'Open', and 'Delete'). Below the object list, there's a search bar for 'Find objects by prefix' and a table for managing object details.

Amazon S3

We're continuing to improve the S3 console to make it faster and easier to use. If you have feedback on the updated experience, please provide feedback.

Provide Feedback

Beckets

Access Points

Object Lambda Access Points

Batch Operations

Access analyzer for S3

Block Public Access settings for this account

Storage Lens

Dashboards

AWS Organizations settings

Feature spotlight

All AWS Marketplace for S3

Amazon S3 > octank-datalake1

octank-datalake1

Objects Properties Permissions Metrics Management Access Points

Objects (1)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

Actions ▾ Create folder Upload

Find objects by prefix

| Name | Type | Last modified | Size | Storage class |
|-------|--------|---------------|------|---------------|
| 2021/ | Folder | - | - | - |

Amazon S3

We're continuing to improve the S3 console to make it faster and easier to use. If you have feedback on the updated experience, choose [Provide feedback](#).

Provide feedback

Amazon S3 > octank-datalake1 > 2021/

2021/

[Copy S3 URL](#)

[Objects](#) [Properties](#)

Objects (2)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 Inventory](#) to get a list of all objects in your bucket. For others to access your objects, you'll need to explicitly grant them permissions. [Learn more](#)

[Actions](#) [Create folder](#) [Upload](#)

Find objects by prefix

| <input type="checkbox"/> | Name | Type | Last modified | Size | Storage class |
|--------------------------|------|--------|---------------|------|---------------|
| <input type="checkbox"/> | o1/ | Folder | - | - | - |
| <input type="checkbox"/> | o2/ | Folder | - | - | - |

Amazon S3 > octank-datalake1 > 2021/ > 08/ > 17/ > 09/

Objects (4)

| Name | Type | Last modified | Size | Storage class |
|---|------|---------------------------------------|--------|---------------|
| octank_deliverystream-1-2021-08-17-09-14-62ec7053-a7e2-4f4e-839d-269d795471 | File | August 17, 2021, 11:18:17 (UTC-02:00) | 8.6 kB | Standard |
| octank_deliverystream-1-2021-08-17-09-14-903efad-0276-4b98-aa7a-c1480e51268 | File | August 17, 2021, 11:21:17 (UTC-02:00) | 8.7 kB | Standard |
| octank_deliverystream-1-2021-08-17-09-23-40-c7572a51-4748-4e84-9c79-97769148264 | File | August 17, 2021, 11:28:42 (UTC-02:00) | 8.6 kB | Standard |
| octank_deliverystream-1-2021-08-17-09-28-41-4f416bd-8184-4570-4734-4b0141793409 | File | August 17, 2021, 11:29:42 (UTC-02:00) | 1.4 kB | Standard |

Feedback [Provide feedback](#)

[Copy 55 URLs](#)

objects, you'll need to explicitly grant them permission. Learn more [? \[x\]](#)

[Upload](#)

< 1 > [@](#)

| | Last modified | Size | Storage class |
|--|---------------------------------------|--------|---------------|
| | August 17, 2021, 11:18:37 (UTC+02:00) | 8.6 KB | Standard |
| | August 17, 2021, 11:21:37 (UTC+02:00) | 8.7 KB | Standard |
| | August 17, 2021, 11:28:42 (UTC+02:00) | 8.6 KB | Standard |
| | August 17, 2021, 11:29:42 (UTC+02:00) | 1.4 KB | Standard |

The screenshot shows the AWS S3 console interface. On the left, there's a sidebar with various AWS services like Amazon S3, AWS Lambda, and CloudWatch Metrics. The main area has a search bar at the top with the query "Athena". Below it, there are two sections: "Services" and "Documentation".

Services

- Athena: Query data in S3 using SQL.

Documentation

- Identity and Access Management In Athena - Amazon Athena [User Guide]
- Resilience in Athena - Amazon Athena [User Guide]
- Troubleshooting in Athena - Amazon Athena [User Guide]
- Monitoring Athena Queries with CloudWatch Events - Amazon Athena [User Guide]

Knowledge Articles

- Resilience in Athena - Amazon Athena [Amazon Athena]
- What Is Amazon Athena? - Amazon Athena [Amazon Athena]
- When should I use Athena? - Amazon Athena [Amazon Athena]

On the right side, there's a preview of a file list with four items:

| Last modified | Size | Storage class |
|---------------------------------------|--------|---------------|
| August 17, 2021, 11:18:37 (UTC+02:00) | 8.6 kB | Standard |
| August 17, 2021, 11:23:17 (UTC+02:00) | 8.7 kB | Standard |
| August 17, 2021, 11:28:42 (UTC+02:00) | 8.6 kB | Standard |
| August 17, 2021, 11:29:42 (UTC+02:00) | 1.4 kB | Standard |

Athena [Query editor](#) [Saved queries](#) [History](#) [Data sources](#) [Workgroup primary](#)

Data source [Connect data source](#)

Amazon Data Catalog

Database

Filter tables and views...

New query 1

Run query Save as Create

Ctrl + Enter to run query, Ctrl + Space to autocomplete

Format query Clear

Athena engine version 2 Release versions (F)

Results

Amazon S3 Services

Feedback

Provide feedback

Copy S3 URL

Search results for 'Athena'

Services (1)

Athena

Query data in S3 using SQL

Documentation (30,971)

See all 30,971 results in Documentation

Identity and Access Management in Athena - Amazon Athena

User Guide

Resilience in Athena - Amazon Athena

User Guide

Troubleshooting in Athena - Amazon Athena

User Guide

Monitoring Athena Queries with CloudWatch Events - Amazon Athena

User Guide

Knowledge Articles (30)

See all 30 results in Knowledge Articles

Resilience in Athena - Amazon Athena

Amazon Athena

What is Amazon Athena? - Amazon Athena

Amazon Athena

When should I use Athena? - Amazon Athena

Amazon Athena

Last modified

| | Size | Storage class |
|---------------------------------------|--------|---------------|
| August 17, 2021, 11:18:37 (UTC-02:00) | 8.6 kB | Standard |
| August 17, 2021, 11:21:37 (UTC-02:00) | 8.7 kB | Standard |
| August 17, 2021, 11:28:42 (UTC-02:00) | 8.6 kB | Standard |
| August 17, 2021, 11:29:42 (UTC-02:00) | 1.4 kB | Standard |

Athena Query editor Saved queries History Data sources Workgroup: primary [Option+1] [Edit](#) [Run](#) [Stop](#) [Cancel](#)

New query 1

1

Data source: AvnDataCataloging Connect data source

Database: Default

Filter tables and views

Tables (0) Create table

The selected database has no tables.

Views (0) Create view

You have not created any views. To create a view, run a query and click "Create view from query".

Run query Save as Create -

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

Athena engine version 2 - Release version: 1.0

Settings Tutorial Help What's new

Formal query Clear

Athena Query editor Saved queries History Data sources Workgroup: primary [Options+I]

Data source: AmazonDataCatalog Connect data source

Database: default

Filter tables and views

Tables (0) Create table

The selected database has no tables.

Views (0) Create view

You have not created any views. To create a view, run a query and click "Create view from query".

New query 1 +

Run query Save as Create -

Use Ctrl + Enter to run query. Ctrl + Space to autocomplete

Format query Clear

Athena Engine version 2 - Release version: 1.0

Results

Connect data source

Step 1: Choose a data source

Choose where your data is located

Athena queries data where it is. Data is not loaded or moved. Learn more [?]

 Query data in Amazon S3

Choose an external data catalog

 Query a data source

Configure a connector for common data sources.



Choose a metadata catalog

The catalog contains the schema for the source data such as column names, data types and table names. Learn more [?]

 AWS Glue Data Catalog Apache Hive metastore

Cancel

Next

Connect data source**Step 1: Choose a data source**

Choose where your data is located

Athena queries data where it is. Data is not loaded or moved. [Learn more](#)  **Query data in Amazon S3**

Choose an external data catalog

 **Query a data source**

Configure a connector for common data sources.



Choose a metadata catalog

The catalog contains the schema for the source data such as column names, data types and table names. [Learn more](#) **AWS Glue Data Catalog****Apache Hive metastore**[Cancel](#)**Next** 

Connect data source

Step 1. Choose a data source

Step 2. Connection details

Connection details: choose an AWS Glue Data Catalog

Choose an AWS Glue Data Catalog in your account or in another account.

 AWS Glue Data Catalog in this account

Create a table in AWS Glue Data Catalog

 AWS Glue Data Catalog in another account

Register an external AWS Glue Data Catalog for cross-account access.

Choose a way to create a table

Athena will connect to your data stored in Amazon S3 and use AWS Glue Data Catalog to store metadata, such as table and column names. Once connected, your databases, tables and views appear in Athena's query editor. [Learn more](#) Create a crawler in AWS Glue

Add a table by setting up a crawler in AWS Glue to analyse data and retrieve the schema automatically

 Create a table using the Athena table wizard

Add a table by entering schema information manually in Athena

[Cancel](#)[Previous](#)[Connect to AWS Glue](#)

- Crawler Info
- Crawler source type
- Data store
- IAM role
- Schedule
- Output
- Previous API results

Add information about your crawler

Crawler name

Enter name

► Tags, description, security configuration, and classifiers (optional)

Next

- Crawler info
- Crawler source type
- Data store
- IAM Role
- Schedule
- Output
- Preview all steps

Add information about your crawler

Crawler name

octane-created

► Tags, description, security configuration, and classifiers (optional)

Next

- Crawler info
 ``stlMark_crawler`
- Crawler source type
- Data store
- IAM Role
- Schedule
- Output
- Review all steps

Specify crawler source type

Choose Existing catalog tables to specify catalog tables as the crawler source. The selected tables specify the data stores to crawl. This option doesn't support JDBC data stores.

Crawler source type

- Data stores
- Existing database tables

Repeat crawls of S3 data stores

- Crawl all folders
Crawl all folders again with every subsequent crawl.
- Crawl new folders only
Only Amazon S3 folders that were added since the last crawl will be crawled. If the schemas are compatible, new partitions will be added to existing tables.

[Back](#)

[Next](#)

- Crawler info
extant_crawler
- Crawler source type
Data stores
- Data store
- ODM logic
- Procedure
- Output
- Review all steps

Add a data store

Choose a data store

53

Connection

Select a connection

Optional. Include a network connection to use with this ED target. Note that each crawler is limited to one network connection so any future ED targets will also use the same connection (or none, if left blank).

Add connection

Crawl data in

Specified path in my account
 Specified path in another account.

Include path

ED:/Document/MyFolder

All folders and files contained in the include path are crawled. For example, type ED:/MySource/MyFolder to crawl all elements in MyFolder within MySource.

Sample size (optional)

Enter an integer between 1 and 249

This field sets the number of files in each result to be returned. If not set, all the files are crawled.

Exclude patterns (optional)

Back Next

- Create info**
- extant_creater
- Create source type**
- Data stores
- Data store
- XML feeds
- Schedules
- Output
- Review all steps

Add a data store

Choose a data store

S3

Connection

Select a connection

Optimally, include a network connection to use with this S3 target. Note that each crawler is limited to one network connection so any future S3 targets will also use the same connection (or none, if left blank).

Add connection

Crawl data in

Specified path in my account
 Specified path in another account.

Include path

S3://DocumentCloud/analyzed

All folders and files contained in the include path are crawled. For example, type S3://MyBucket/MyFolder to crawl all elements in MyFolder within MyBucket.

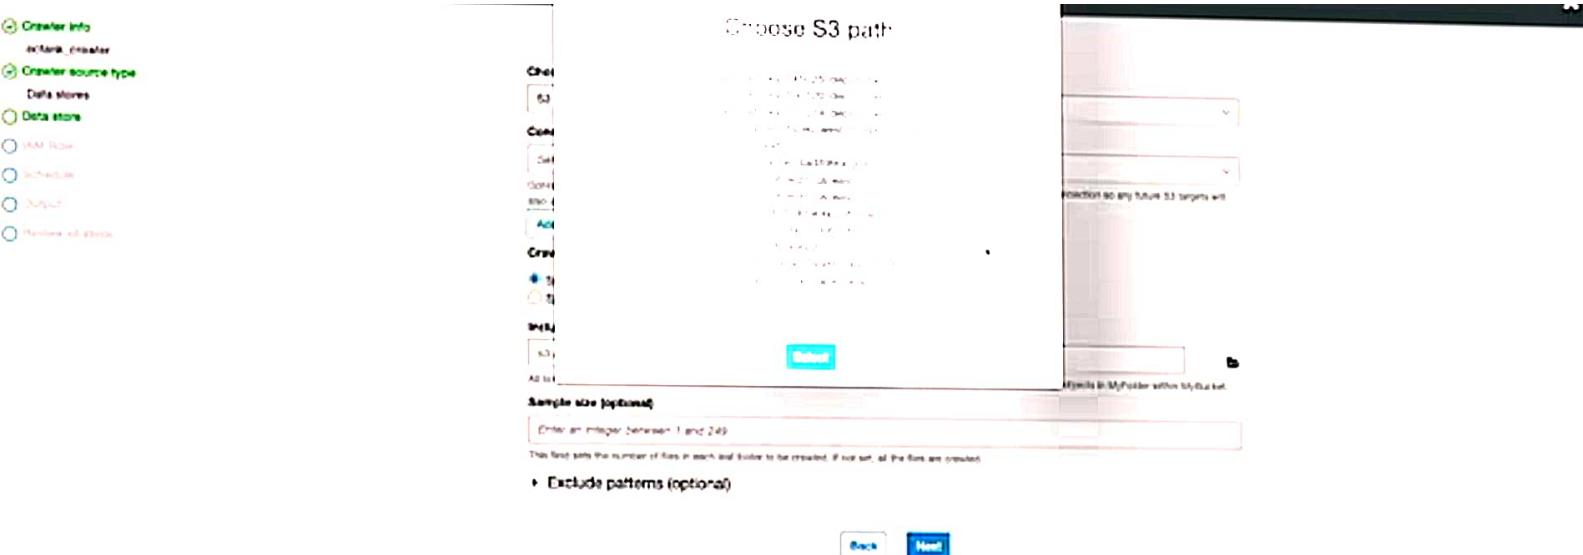
Sample size (optional)

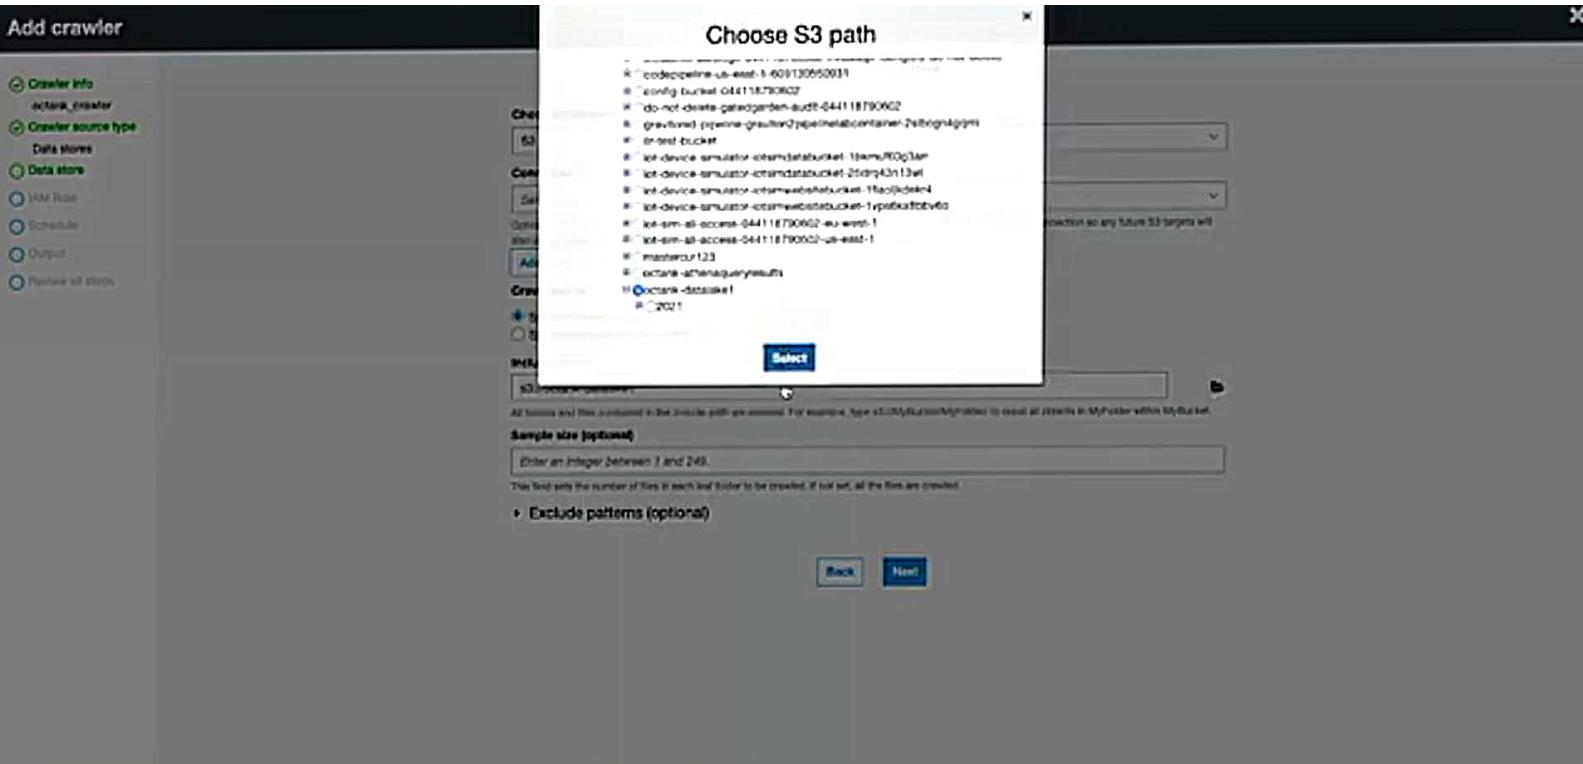
Enter an integer between 1 and 249

This field sets the number of files in each and folder to be crawled. If not set, all the files are crawled.

Exclude patterns (optional)

Back Next





- Crawler info**
extank_crawler
- Crawler source type**
Data stores
- Data store
S3: a3/cocktails-datalake
- IAM roles
- Schedule
- Output
- Preview all steps

Add another data store

Yes
 No

[Back](#) [Next](#)

Chosen data stores
S3: a3/cocktails-datalake

- Crawler info**
extank_crawler
- Crawler source type**
Data stores
- Data store**
S3: s3://octane-database1
- IAM Role**
- S3 module
- Output
- Review all stages

Choose an IAM role

The IAM role allows the crawler to run and access your Amazon S3 data stores. [Learn more](#)

- Update a policy in an IAM role
- Choose an existing IAM role
- Create an IAM role

IAM role 

AWSGlueServiceRole-extane_datalake

To create an IAM role, you must have `CreateRole`, `CreatePolicy`, and `AttachRolePolicy` permissions.

Create an IAM role named '`AWSGlueServiceRole-torname`' and attach the AWS managed policy `AWSGlueServiceRole`, plus an inline policy that allows read access to:

- S3:octane-database1

You can also create an IAM role on the [IAM console](#).

[Back](#)

[Next !\[\]\(79e06214067bdb5b59865fb8fd052c24_img.jpg\)](#)

- Crawler info**
octark_crawler
- Crawler source type**
Data stores
- Data store**
S3 s3://octark-data...
Amazon S3
- IAM Role**
arn:aws:iam::6441187:role/ServiceRoleForAWSGlueService
Role-
octark_databricks
- Schedule**
- Output**
- Review all steps**

Create a schedule for this crawler

Frequency

Run on demand

Back

Next

- Crawler info**
extarnl_crawler
- Crawler source type**
Data stores
- Data store**
S3: s3://Portaria-del...
- IAM Role**
arn:aws:iam::6441187:80602role/service-role/AmazonSGlueServiceRole-
stark_databrewrole
- Schedule**
Run on demand
- Output**
- Review all steps**

Configure the crawler's output

Database 0

Choose a database to contain tables

Add database

Prefix added to tables (optional) 0

Type a prefix added to table names

► Grouping behavior for S3 data (optional)

► Configuration options (optional)

Back

Next

- Crawler info**
octank_crawler
- Crawler source type**
Data stores
- Data store**
S3: s3://octank-datalake
- IAM Role**
arn:aws:iam::0441187:
90602:role/service-
role-AllSGGiveService
Role-
octank_datalakerole
- Schedule**
Run on demand
- Output**
- Review all steps**

Configure the crawler's output

Database 0

Choose a database to contain tables

Add database 

Prefix added to tables (optional) 

Type a prefix added to table names

▶ Grouping behavior for S3 data (optional)

▶ Configuration options (optional)



Add database

Database name

Description and location (optional)

Resource link name

Shared database suggestions

Shared database

Shared database owner account ID





Add database

Database name

Description and location (optional)

Resource link name

Shared database suggestions

Shared database

Shared database owner account ID

Create

https://eu-west-1.console.aws.amazon.com/glue/home?region=eu-west-1#addCrawler

Add crawler

Crawler info
octank_crawler

Crawler source type
Data stores

Data store
S3: s3://octank-database

IAM Role
arn:aws:iam::04411879002:role/service-role/AWSGlueServiceRole-octank_database

Schedule
Run on demand

Output
octank_db

Review all steps

Crawler info

Name: octank_crawler
Tags: -

Data stores

Data store: S3
Include path: s3://octank-database/
Connection:
Exclude patterns: -

IAM role

IAM role: arn:aws:iam::04411879002:role/service-role/AWSGlueServiceRole-octank_database

Schedule

Schedule: Run on demand

Output

Database: octank_db
Prefix added to tables (optional):
Create a single schema for each S3 path: false
Table level (optional):
• Configuration options

Back Finish

https://eu-west-1.console.aws.amazon.com/glue/home?region=eu-west-1#AddCrawler

Add crawler

Crawler info
octank_crawler

Crawler source type
Data stores

Data store
S3: s3://octank-database

IAM Role
arn:aws:iam::044118790602:role/service-role/AmazonSQSRole

Schedule
Run on demand

Output
octank_db

Review all steps

Crawler Info

Name: octank_crawler
Tags: -

Data stores

Data store: S3
Include path: s3://octank-database/
Connection:
Exclude patterns:

IAM role

IAM role: arn:aws:iam::044118790602:role/service-role/AmazonSQSRole

Schedule

Schedule: Run on demand

Output

Database: octank_db
Prefix added to tables (optional):
Create a single schema for each S3 path: false
Table level (optional):
• Configuration options

[Back](#) [Finish](#)

AWS Glue

Data catalog

[Databases](#)[Tables](#)[Connections](#)[Crawlers](#)[Classifiers](#)[Schema registries](#)[Schemas](#)[Settings](#)

ETL

[AWS Glue Studio](#)[Blueprints](#)[Workflows](#)[Jobs](#)[ML Transforms](#)[Triggers](#)[Dev endpoints](#)[Notebooks](#)

Security

[Security configurations](#)

Tutorials

[Add crawler](#)[Explore table](#)[Add job](#)[Resources](#)[What's new](#)

Crawlers - A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

[User preferences](#)

| <input type="checkbox"/> | Name | Schedule | Status | Logs | Last runtime | Median runtime | Tables updated | Tables added |
|--------------------------|----------------|----------|--------|------|--------------|----------------|----------------|--------------|
| <input type="checkbox"/> | octank_crawler | | Ready | | 0 secs | 0 secs | 0 | 0 |

AWS Glue

Data catalog

Databases

Tables

Connections

Crawlers

Classifiers

Schema registries

Schemas

Settings

ETL

AWS Glue Studio

Blueprints

Workflows

Jobs

ML Transforms

Triggers

Dev endpoints

Notebooks

Security

Security configurations

Tutorials

Add crawler

Explore table

Add job

Resources

DF

What's new

Crawlers - A crawler connects to a data store, progresses through a prioritized list of classifiers to determine the schema for your data, and then creates metadata tables in your data catalog.

User preferences

| <input type="checkbox"/> | Name | Schedule | Status | Logs | Last runtime | Median runtime | Tables updated | Tables added |
|--------------------------|----------------|----------|----------|------|--------------|----------------|----------------|--------------|
| <input type="checkbox"/> | octane_crawler | | Starting | | 0 secs | 0 secs | 0 | 0 |

https://eu-west-1.console.aws.amazon.com/athena/home?#workset=watch&tab=query

Athena Query editor Saved queries History Data sources Workgroup: primary Options [1]

Data source: Connect data source

New query 1 +

1

Database:

- default
- default
- athena_ls
- sampled
- test

Views (0) Create view

You have not created any views. To create a view, run a query and click "Create view from query".

Run query Save as Create

Ctrl + Enter to run query, Ctrl + Space to autocomplete

Format query Clear

Athena engine version 2 - Release versions [?]

Results

This screenshot shows the Amazon Athena Query Editor. On the left, there's a sidebar with 'Data sources' (Athena Data Catalog), 'Database' (with options like default, athena_ls, sampled, test), and 'Views (0)'. The main area has a 'New query 1' tab open. It contains a single row with the number '1'. Below the query area are buttons for 'Run query', 'Save as', and 'Create'. A note says 'You have not created any views. To create a view, run a query and click "Create view from query".' At the bottom, there are 'Format query' and 'Clear' buttons, and a note about the Athena engine version.

Athena Query editor Saved queries History Data sources Workgroup: primary

Data source: AmazonDataCatalog Connect data source

Database: octane_db Connect database

Tables (1) Create table

- octane_database1 (Partitioned)

Views (0) Create view

You have not created any views. To create a view, run a query and click "Create view from query".

New query 1

Run query Save as Create

Use Ctrl + Enter to run query. CM + Scroll to autocomplete

Format query Clear

Athena engine version 2 - Release version 1

Results

The screenshot shows the AWS Athena Query Editor. On the left, there's a sidebar with navigation links like 'Athena', 'Query editor', 'Saved queries', 'History', 'Data sources', and 'Workgroup: primary'. Under 'Data sources', 'AmazonDataCatalog' is selected. Under 'Database', 'octane_db' is selected. A table named 'octane_database1 (Partitioned)' is listed under 'Tables (1)'. Below it, a note says 'You have not created any views. To create a view, run a query and click "Create view from query".' In the main area, a new query window titled 'New query 1' is open. It contains a text input field with placeholder text 'Use Ctrl + Enter to run query. CM + Scroll to autocomplete.', and several buttons: 'Run query', 'Save as', 'Create', 'Format query', and 'Clear'. At the bottom right, it says 'Athena engine version 2 - Release version 1'.

Athena Query editor Saved queries History Data sources Workgroup: primary

Data source: AvnDataCatalog Connect data source

Database: extana_db

Filter tables and views

Tables (1) Create table: [extana_database1 \(Partitioned\)](#)

Views (0) Create view

You have not created any views. To create a view, run a query and click "Create view from query".

New query 1

Run query Save as Create

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

Athena engine version 2 - Release version 1

The screenshot shows the AWS Athena Query Editor. On the left, there's a sidebar with 'Data source' set to 'AvnDataCatalog' and 'Database' set to 'extana_db'. Below this, sections for 'Tables (1)' and 'Views (0)' are shown, with a note about creating views from queries. The main area is titled 'New query 1' and contains buttons for 'Run query', 'Save as', and 'Create'. At the bottom, it says 'Use Ctrl + Enter to run query, Ctrl + Space to autocomplete'. The top navigation bar includes tabs for 'Query editor', 'Saved queries', 'History', 'Data sources', and 'Workgroup: primary'. The top right has links for 'Settings', 'Tutorial', 'Help', and 'What's new'.

Athena Query editor Saved queries History Data sources Workgroup: primary

Data source: AmazonDataCatalog Connect data source

Database: octane_db

Tables (1): octane_database1 (Partitioned)

Views (0):

You have not created any views. To create a view, run a query and click "Create view from query".

New query 1

Preview table Show properties Delete table List partitions Generate Create Table DDL

Run query Save as Create -

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

Athena Engine version 2 - Release version: 0

The screenshot shows the AWS Athena Query Editor interface. At the top, there's a navigation bar with tabs for 'Athena', 'Query editor' (which is selected), 'Saved queries', 'History', 'Data sources', and 'Workgroup: primary'. Below the navigation bar, there are sections for 'Data source' (set to 'AmazonDataCatalog') and 'Database' (set to 'octane_db'). On the left, there's a sidebar with sections for 'Tables (1)' (listing 'octane_database1 (Partitioned)') and 'Views (0)'. A note below the sidebar says, 'You have not created any views. To create a view, run a query and click "Create view from query".' In the main workspace, a new query is being created with the name 'New query 1'. A context menu is open over the table 'octane_database1', showing options like 'Preview table', 'Show properties', 'Delete table', 'List partitions', and 'Generate Create Table DDL'. Below the table list, there are buttons for 'Run query', 'Save as', and 'Create -'. A note at the bottom of the workspace says, 'Use Ctrl + Enter to run query, Ctrl + Space to autocomplete'. At the very bottom, it says 'Athena Engine version 2 - Release version: 0'.

AWS Services ▾

Amazon Athena Query editor

Saved queries History Data sources Workgroup: primary

Search for services, features, marketplace products, and docs [Options+]

New query 1 New query 2 +

```
1 SELECT * FROM "rectask_db"."rectask_detailed" limit 100
```

Data source Connect data source

AmazonDataCatalog

Database rectask_db

Tables (1) Create table

+ rectask_dataset1 (Partitioned)

Views (0) Create view

You have not created any views. To create a view, run a query and click "Create view from query".

Cancel Save as Create

Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Results

Billing Query ...

Estimated time elapsed: 0 seconds

You can run another query by clicking on the New Query button. The current query will continue to run in the background. You can check the status of all queries in the History tab.

Format query Clear

Athena engine version 2 Remote versions (0)

https://eu-west-1.console.aws.amazon.com/athena/home?#®ion=eu-west-1&queryHistoryId=47c6b-3f46-414a-a080-ccf3b4cb60b1

Athena Query editor Saved queries History Data sources Workgroup: primary Settings Tutorial Help What's new

Data source: AaaDataCatalog Connect data source

Database: octank_db

New query 1 New query 2 +
 : SELECT * FROM "octank_db"."octank_database1" LIMIT 100

Tables (1): Create table
 octank_database1 (Partitioned)

Views (0): Create view

You have not created any views. To create a view, run a query and click "Create view from query".

Run query Save as Create (Run time: 0.72 seconds, Data scanned: 30.29 KB) Format query Clear

Use Ctrl + Enter to run query, CM + Space to autocomplete

Athena engine version 2 - Release version 0.8

Results

| # | deviceid | Timestamp | battery | speed | location | partition_0 | partition_1 | partition_2 | partition_3 |
|----|-----------|------------|---------|-------|--|-------------|-------------|-------------|-------------|
| 1 | encocet1 | 1620308524 | 85 | 42 | [lat=48.11408829432654, long=-11.7020835488034] | 2021 | 05 | 06 | 14 |
| 2 | encocet25 | 1620308529 | 85 | 42 | [lat=48.113511102384474, long=-11.67423185487345] | 2021 | 05 | 06 | 14 |
| 3 | encocet27 | 1620308534 | 76 | 1 | [lat=48.17521915776688, long=-11.66589025927934] | 2021 | 05 | 06 | 14 |
| 4 | encocet19 | 1620308539 | 23 | 96 | [lat=48.17628810041454, long=-11.498270210946587] | 2021 | 05 | 06 | 14 |
| 5 | encocet8 | 1620308544 | 16 | 20 | [lat=48.147213958175835, long=-11.836380808580911] | 2021 | 05 | 06 | 14 |
| 6 | encocet12 | 1620308549 | 83 | 6 | [lat=48.03587783580419, long=-11.534354181808229] | 2021 | 05 | 06 | 14 |
| 7 | encocet22 | 1620308554 | 95 | 18 | [lat=48.2362973536684, long=-11.475849363388267] | 2021 | 05 | 06 | 14 |
| 8 | encocet11 | 1620308559 | 18 | 71 | [lat=48.1216395767278, long=-11.87389173381471] | 2021 | 05 | 06 | 14 |
| 9 | encocet30 | 1620308564 | 17 | 67 | [lat=48.03827717831533, long=-11.731347888521159] | 2021 | 05 | 06 | 14 |
| 10 | encocet10 | 1620308569 | 79 | 25 | [lat=48.16745490290829, long=-11.441915249161589] | 2021 | 05 | 06 | 14 |

https://eu-west-1.console.aws.amazon.com/athena/home?#region=eu-west-1&queryHistoryId=0947ccb-3f46-414a-a080-ccf304ccb6001

Athena Query editor Saved queries History Data sources Workgroup: primary Settings Tutorial Help What's new

Data source: AzaDataCatalog Connect data source

Database: octane_db

Tables (1): octane_database1 (Partitioned) Create table

Views (0): Create view

You have not created any views. To create a view, run a query and click "Create view from query".

New query 1 New query 2 +
 : SELECT * FROM "octane_db"."octane_database1" LIMIT 100

Run query Save as Create (Run time: 0.72 seconds, Data scanned: 30.23 KB)
 Use Ctrl + Enter to run query, Ctrl + Space to autocomplete

Athena engine version 2 - Release version 0.8

Results

| # | deviceid | Timestamp | battery | speed | location | partition_0 | partition_1 | partition_2 | partition_3 |
|----|------------|------------|---------|-------|---|-------------|-------------|-------------|-------------|
| 1 | escooter1 | 1625308524 | 85 | 42 | [lat=48.19408829432654, long=11.7000353488534] | 2021 | 05 | 06 | 14 |
| 2 | escooter25 | 1625308529 | 85 | 42 | [lat=48.113311302384474, long=11.674231854873845] | 2021 | 05 | 06 | 14 |
| 3 | escooter27 | 1625308534 | 78 | 1 | [lat=48.17521915779668, long=11.64589925927934] | 2021 | 05 | 06 | 14 |
| 4 | escooter19 | 1625308539 | 23 | 56 | [lat=48.17808510041454, long=11.489270216986587] | 2021 | 05 | 06 | 14 |
| 5 | escooter68 | 1625308544 | 16 | 20 | [lat=48.147213058175835, long=11.826360087880911] | 2021 | 05 | 06 | 14 |
| 6 | escooter12 | 1625308549 | 83 | 6 | [lat=48.03597783580415, long=11.534354181808029] | 2021 | 05 | 06 | 14 |
| 7 | escooter22 | 1625308554 | 95 | 18 | [lat=48.2362973536884, long=11.43584363886267] | 2021 | 05 | 06 | 14 |
| 8 | escooter11 | 1625308559 | 14 | 71 | [lat=48.1210439937278, long=11.67369113581471] | 2021 | 05 | 06 | 14 |
| 9 | escooter30 | 1625308564 | 11 | 67 | [lat=48.03827118931533, long=11.72134788501159] | 2021 | 05 | 06 | 14 |
| 10 | escooter30 | 1625308569 | 79 | 25 | [lat=48.10745490290829, long=11.44195249161589] | 2021 | 05 | 06 | 14 |

https://eu-west-1.console.aws.amazon.com/athena/home?#queryEditor

QuickSight Connect workflow

Search results for 'Quick'

Services (6)

See all 6 results ▾

- QuickSight**
Fast, easy-to-use business analytics
- AWS Glue DataBrew**
Visual data preparation tool to clean and normalize data for analytics and machine learning
- Amazon Machine Learning**
Build Smart Applications Quickly and Easily
- CodeStar**
Quickly develop, build, and deploy applications

Features (4)

See all 6 results ▾

- Quick Setup**
Systems Manager feature
- Quick start**
IoT Analytics feature
- Quick Connect workflow**
FreeRTOS Native
- Dashboard**
IoT feature

Documentation (20,108)

See all 20,108 results in Documentation ▾

Format query Clear

Athena engine version 2 · Release version 1.0

| | partition_0 | partition_1 | partition_2 | partition_3 |
|--------|-------------|-------------|-------------|-------------|
| 10034 | 2021 | 05 | 06 | 14 |
| 12345 | 2021 | 05 | 06 | 14 |
| 21934 | 2021 | 05 | 06 | 14 |
| 46887 | 2021 | 05 | 06 | 14 |
| 540111 | 2021 | 05 | 06 | 14 |
| 60629 | 2021 | 05 | 06 | 14 |
| 64267 | 2021 | 05 | 06 | 14 |
| 71413 | 2021 | 05 | 06 | 14 |
| 71158 | 2021 | 05 | 06 | 14 |
| 81589 | 2021 | 05 | 06 | 14 |

https://eu-west-1.console.aws.amazon.com/athena/home?®ion=eu-west-1&queryHistoryID=47cbdb-3546-414a-a080-ccf3b4cb0d81

Search results for 'quick'

Services (6)

Features (4)

Documentation (20,104)

Knowledge Articles (30)

Marketplace (3,491)

Services

See all 6 results ▾

- QuickSight**
Fast, easy-to-use business analytics
- AWS Glue DataBrew**
Visual data preparation tool to clean and normalize data for analytics and machine learning
- Amazon Machine Learning**
Build Smart Applications Quickly and Easily
- CodeStar**
Quickly develop, build, and deploy applications

Features

See all 6 results ▾

- Quick Setup**
Systems Manager feature
- Quick start**
IoT Analytics feature
- Quick Connect workflow**
FreeRTOS feature
- Dashboard**
EC2 feature

Documentation

See all 20,104 results in Documentation ▾

partition_0 + partition_1 + partition_2 + partition_3 +

| | | | | |
|--------|------|----|----|----|
| 15634 | 2021 | 05 | 06 | 14 |
| 17345 | 2021 | 05 | 06 | 14 |
| 27934 | 2021 | 05 | 06 | 14 |
| 46587 | 2021 | 05 | 06 | 14 |
| 580111 | 2021 | 05 | 06 | 14 |
| 580129 | 2021 | 05 | 06 | 14 |
| 58267 | 2021 | 05 | 06 | 14 |
| 61410 | 2021 | 05 | 06 | 14 |
| 71158 | 2021 | 05 | 06 | 14 |
| 91589 | 2021 | 05 | 06 | 14 |

Format query Clear

Athena engine version 2 - Release versions (5)

Find analyses & more

Datasets

New dataset

★ Favorites

Recent

My folders

Shared folders

Dashboards

Analyses

Datasets

No datasets

Import or create a new dataset to start an analysis.

QuickSight

Datasets

Create a Dataset
FROM NEW DATA SOURCES

Upload a file (CSV, Excel, JSON, XML, JDBC, JAR) Salesforce Connect to Salesforce

Athena RDS Redshift Aurora MySQL PostgreSQL Oracle SQL Server Aurora MySQL PostgreSQL AWS IoT Analytics Amazon Elasticsearch Service Spark Teradata Snowflake Jira ServiceNow Timestream GitHub Twitter Adobe Analytics

New Athena data source

Data source name: Enter a name for the data source.

Athena workgroup: [Empty]

Validate connection: SSL is enabled

Create data source

SPICE capacity for this region: 28 SPICE of 2250

Admin/Deployment Manager

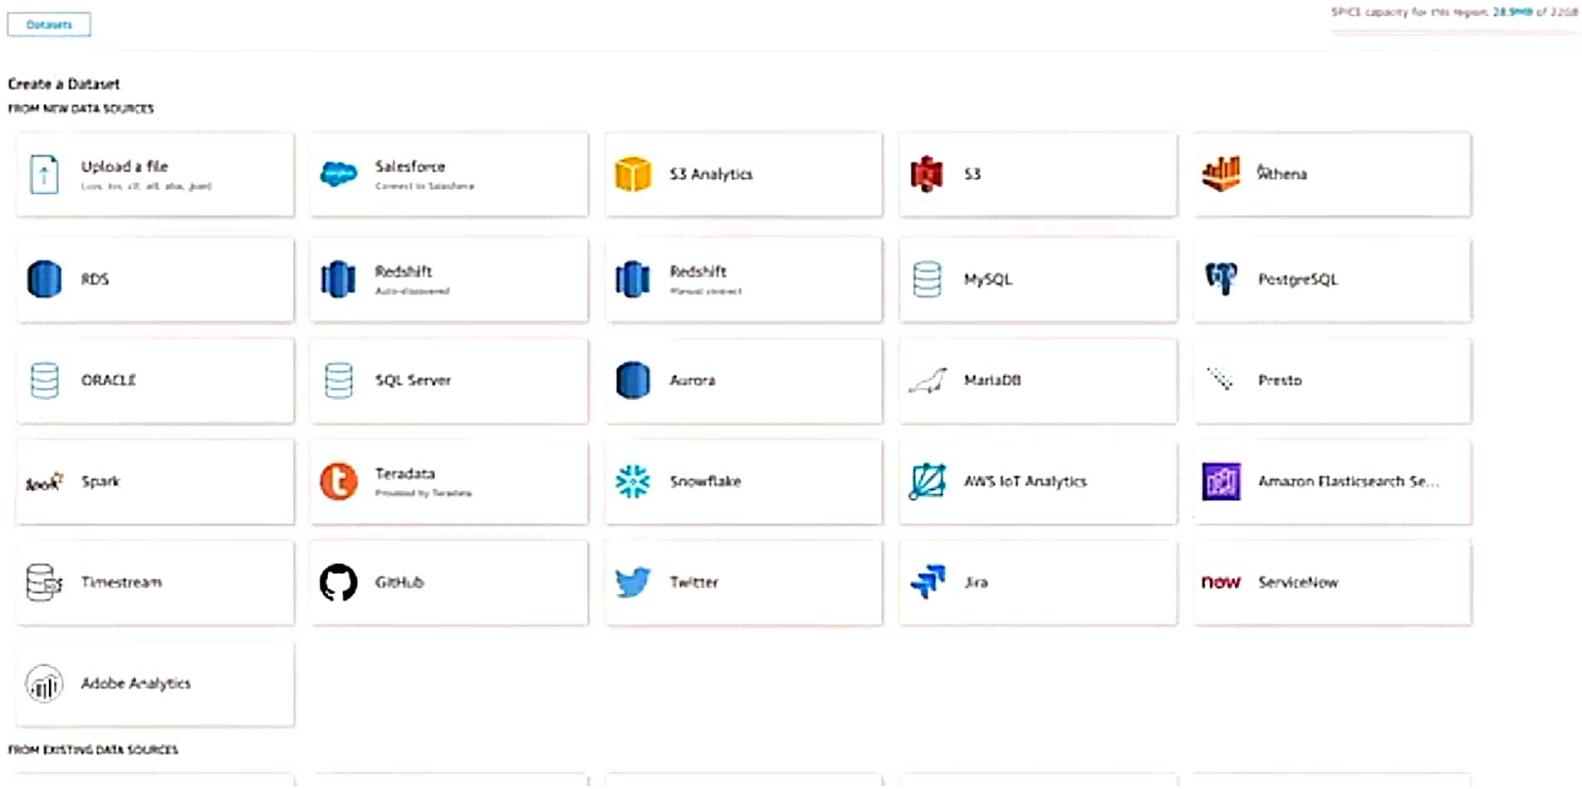
This screenshot shows the 'Data Sources' section of the Amazon QuickSight console. A modal window titled 'New Athena data source' is open in the center. It contains fields for 'Data source name' (with placeholder 'Enter a name for the data source.'), 'Athena workgroup' (set to '[Empty]'), and connection validation status ('Validate connection' - 'SSL is enabled'). At the bottom right of the modal is a blue 'Create data source' button. The background shows a grid of other data source icons, including RDS, Redshift, Aurora, MySQL, PostgreSQL, Oracle, SQL Server, Aurora MySQL, PostgreSQL, AWS IoT Analytics, Amazon Elasticsearch Service, Spark, Teradata, Snowflake, Jira, ServiceNow, Timestream, GitHub, Twitter, and Adobe Analytics. The top right corner of the screen displays 'SPICE capacity for this region: 28 SPICE of 2250' and 'Admin/Deployment Manager'.

Datasets

Create a Dataset

FROM NEW DATA SOURCES

| | | | | |
|--|---|---|---|--|
|  Upload a file <small>(CSV, AVRO, JSON, XML, Parquet, JDBC)</small> |  Salesforce <small>Connect to Salesforce</small> |  S3 Analytics |  S3 |  Athena |
|  RDS |  Redshift <small>Auto-discovery</small> |  Redshift <small>Manual connect</small> |  MySQL |  PostgreSQL |
|  ORACLE |  SQL Server |  Aurora |  MariaDB |  Presto |
|  Spark |  Teradata <small>Provided by Teradata</small> |  Snowflake |  AWS IoT Analytics |  Amazon Elasticsearch Se... |
|  Timestream |  GitHub |  Twitter |  Jira |  ServiceNow |



QuickSight

Datasets

Create a Dataset
FROM NEW DATA SOURCES

Upload a file (CSV, XLS, XML, JSON)

Salesforce Connect to Salesforce

Athena

RDS Redshift MySQL PostgreSQL

ORACLE SQL Server Aurora MariaDB Presto

Spark Teradata Snowflake AWS IoT Analytics Amazon Elasticsearch Se...

Timestream GitHub Twitter Jira ServiceNow

Adobe Analytics

FROM EXISTING DATA SOURCES

New Athena data source

Data source name: octank_datasource

Athena workgroup: [empty]

Validate connection: SSL is enabled

Create data source

Datasets

Create a Dataset

FROM NEW DATA SOURCES

- Upload a file (CSV, JSON, XML, XML, JSON)
- Salesforce Connect to Salesforce
- RDS
- ORACLE
- Spark
- Timestream
- Adobe Analytics

FROM EXISTING DATA SOURCES

Choose your table

octank_databoar

Catalog: contain sets of databases.

AwsDataCatalog

Database: contain sets of tables.

Select...

- default
- octank_db
- sampledb
- test

Edit/Preview Data Use custom SQL Select

SPICE capacity for this region: 28.9998 of 2268

Athena

PostgreSQL

Presto

AWS IoT Analytics

Amazon Elasticsearch Se...

Jira

ServiceNow

The screenshot shows the AWS Glue console interface for creating a new dataset. On the left, there's a grid of icons for various data sources like RDS, Oracle, and Spark. A modal window titled 'Choose your table' is open in the center. It displays the database name 'octank_databoar'. Below it, there are two dropdown menus: 'Catalog: contain sets of databases.' set to 'AwsDataCatalog' and 'Database: contain sets of tables.' set to 'Select...'. The 'Select...' dropdown contains four options: 'default', 'octank_db', 'sampledb', and 'test'. At the bottom of the modal are three buttons: 'Edit/Preview Data', 'Use custom SQL', and a blue 'Select' button. To the right of the modal, there's a summary of SPICE capacity: 'SPICE capacity for this region: 28.9998 of 2268'. Below the modal, there's another grid of icons for services like Athena, PostgreSQL, Presto, AWS IoT Analytics, Amazon Elasticsearch Service, Jira, and ServiceNow.

Datasets

Create a Dataset
FROM NEW DATA SOURCES

Upload a file (CSV, AVRO, JSON, XML, Parquet)
Salesforce Connect to Salesforce

RDS Redshift Auto-discovered

ORACLE SQL Server

Spark Teradata Provided by Teradata

Timestream GitHub Twitter

Adobe Analytics

Choose your table
ectank_datadome

Catalog: contain sets of databases.
AmazonDataCatalog

Database: contain sets of tables.
ectank_0ff

Table: contain the data you can visualize.
 ectank_datadome1

View previous data Use custom SQL

SPICE capacity for this region: 28.99GB of 22GB

Athena PostgreSQL Presto AWS IoT Analytics Amazon Elasticsearch Se... Jira ServiceNow

Datasets

Create a Dataset
FROM NEW DATA SOURCES

Upload a file (CSV, XLS, XML, JSON, PDF)
Salesforce Connect to Salesforce

RDS Redshift Auto-discovered

ORACLE SQL Server

Spark Teradata Provided by Teradata

Timestream GitHub Twitter

Adobe Analytics

Choose your table ectank_datasource

Catalog: contain sets of databases.
AmazonDataCatalog

Database: contain sets of tables.
ectank_db

Table: contain the data you can visualize.
ectank_datasource

Edit/Preview data Use custom SQL Select

Athena PostgreSQL Presto AWS IoT Analytics Amazon Elasticsearch Service Jira ServiceNow

SPICE capacity for this region: 28.99GB of 22GB

Datasets

Create a Dataset
FROM NEW DATA SOURCES

Upload a file
(CSV, AVRO, JSON, XML, Parquet, PDF)

Salesforce
Connect to Salesforce

RDS

ORACLE

Spark

Timestream

Adobe Analytics

FROM EXISTING DATA SOURCES

Athena

PostgreSQL

Presto

AWS IoT Analytics

Amazon Elasticsearch Se...

Jira

ServiceNow

Choose your table
octank_datalake

Catalog: contain sets of databases.
AvnDataCatalog

Database: contain sets of tables.
octank_db

Table: contain the data you can visualize.
 octank_datalake?

Edit/Preview data Use custom SQL Select

SPICE capacity for this region: 28.948 of 2200

The screenshot shows the AWS Data Studio interface. On the left, there's a grid of data source icons under 'FROM NEW DATA SOURCES' and 'FROM EXISTING DATA SOURCES'. A modal window titled 'Choose your table' is open in the center. It displays the catalog 'AvnDataCatalog', the database 'octank_db', and the table 'octank_datalake?'. There are buttons for 'Edit/Preview data' and 'Use custom SQL', and a prominent blue 'Select' button at the bottom right of the modal.

Finish dataset creation

Table: octank_datalake1
Data source: octank_datalake1
Schema: octank_db

Import to SPICE for quicker analytics ✓ 12GB available

Directly query your data

Email owners when a refresh fails

[Edit/Preview data](#) [Augment with SageMaker](#) [Visualize](#)

SPICE capacity for this region: 28.00GB of 22GB

Datasets

Create a Dataset FROM NEW DATA SOURCES

Upload a file (CSV, JSON, XML, AVRO, Parquet) Salesforce Connect to Salesforce

RDS Redshift Auto-discovered MySQL Athena

ORACLE PostgreSQL

Apache Spark MariaDB Presto

Timestream Amazon Elasticsearch Se...

GitHub Snowflake AWS IoT Analytics

Adobe Analytics Twitter Jira ServiceNow

FROM EXISTING DATA SOURCES

SAPCE capacity for this region: 28.59% of 7,200

Create a Dataset
FROM NEW DATA SOURCES

Upload a file (CSV, JSON, XML, AVRO, Parquet, JSONL)

Salesforce Connect to Salesforce

RDS Auto-discovered

ORACLE

Spark

Timestream

Adobe Analytics

Redshift

SQL Server

Teradata Provided by Teradata

Github

Twitter

Aurora

Snowflake

MySQL

MariaDB

AWS IoT Analytics

Jira

Athena

PostgreSQL

Presto

Amazon Elasticsearch Service

ServiceNow

Finish dataset creation

Name: customer-database

Data source: Redshift

Region: us-east-1

Request to AWS Lambda for customer analysis

0 of 1250 available

Dataset quality: good

CloudWatch Metrics for refresh tasks

Edit preview data Segment with Segmenter

Preview

The screenshot shows the Amazon QuickSight interface for creating visualizations. On the left, there's a sidebar with various icons for navigation and settings. The main area is titled "AutoGraph" and contains the message: "Choose 1 or more fields and let QuickSight choose the most appropriate chart". At the top left, the dataset "details_dataset1" is selected. The top right shows an "Importing:" status box indicating "0 rows were imported to SPICE" and "0 rows were skipped".

Dataset: details_dataset1

Field wells:

Sheet 1 +

Importing:

0 rows were imported to SPICE
0 rows were skipped

Options

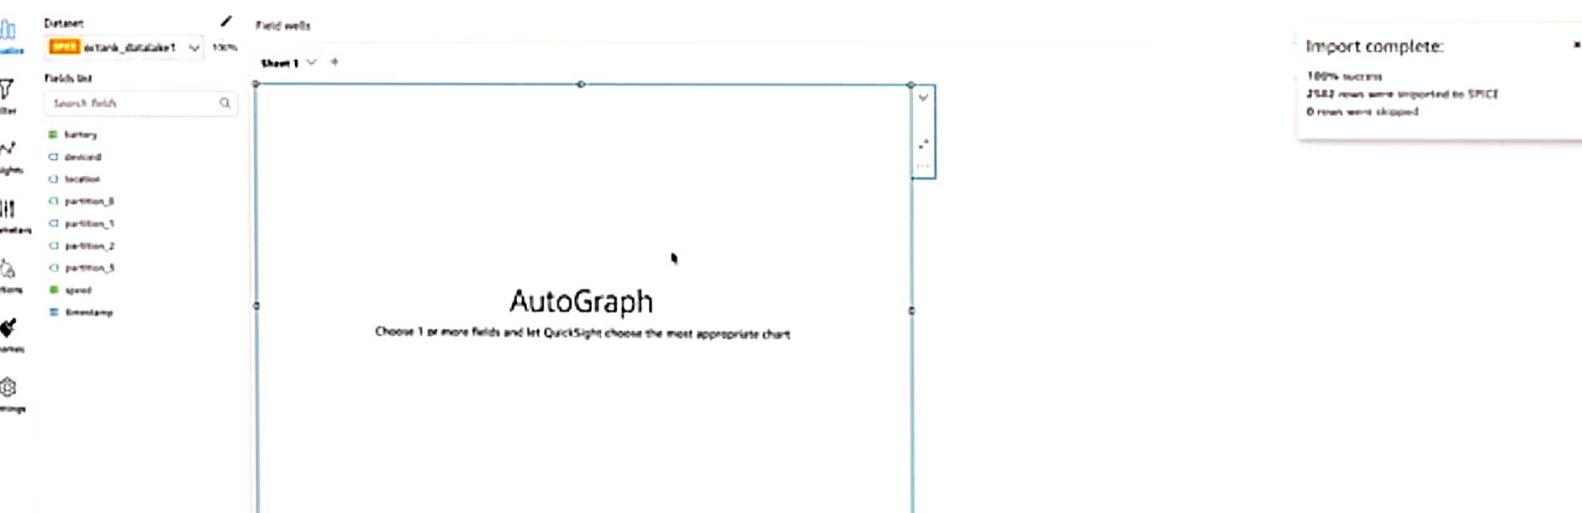
Fields list

Search field:

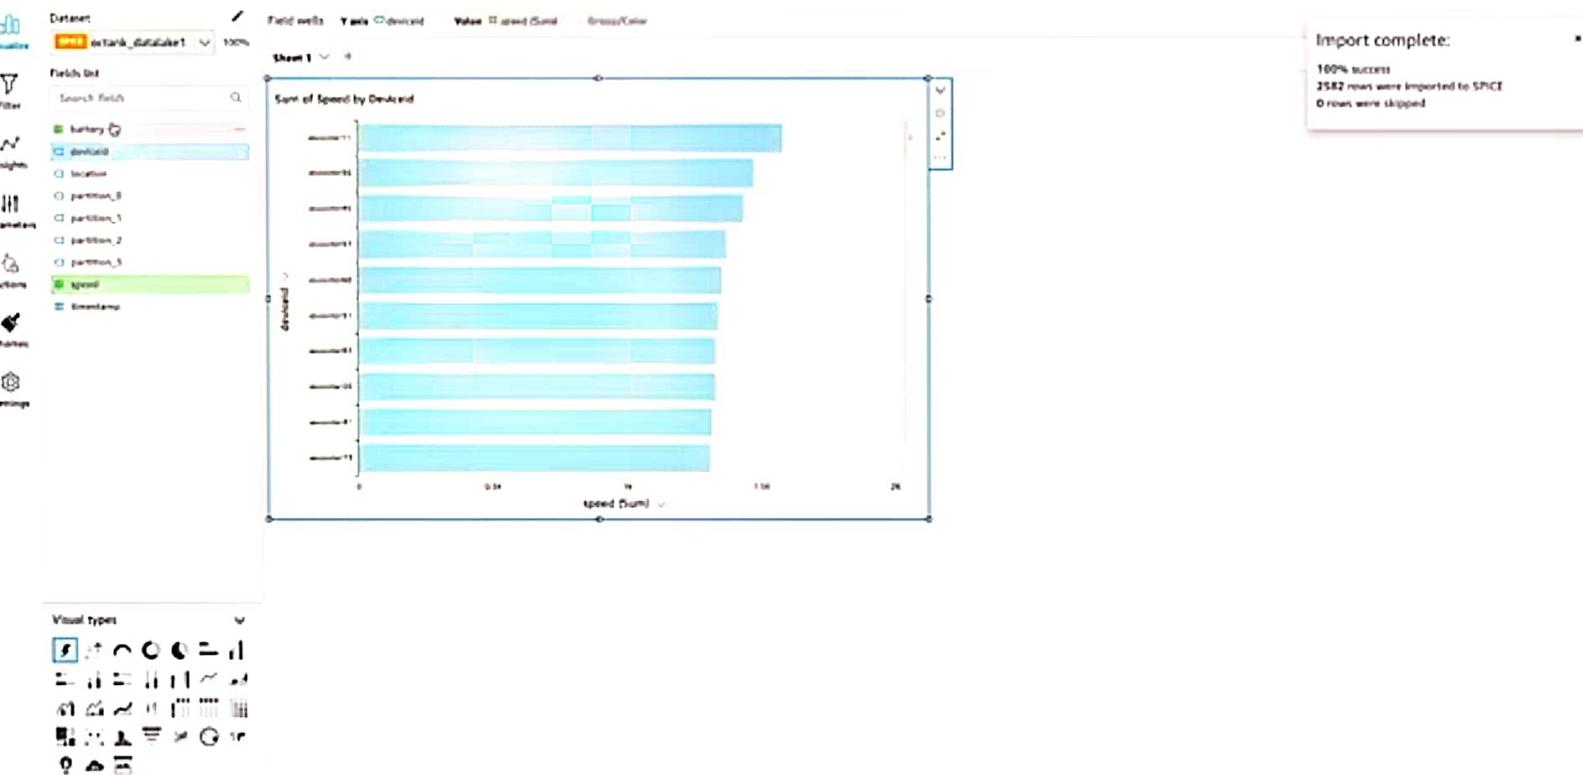
- battery
- deviceid
- location
- partition_0
- partition_1
- partition_2
- partition_3
- speed
- timestamp

AutoGraph

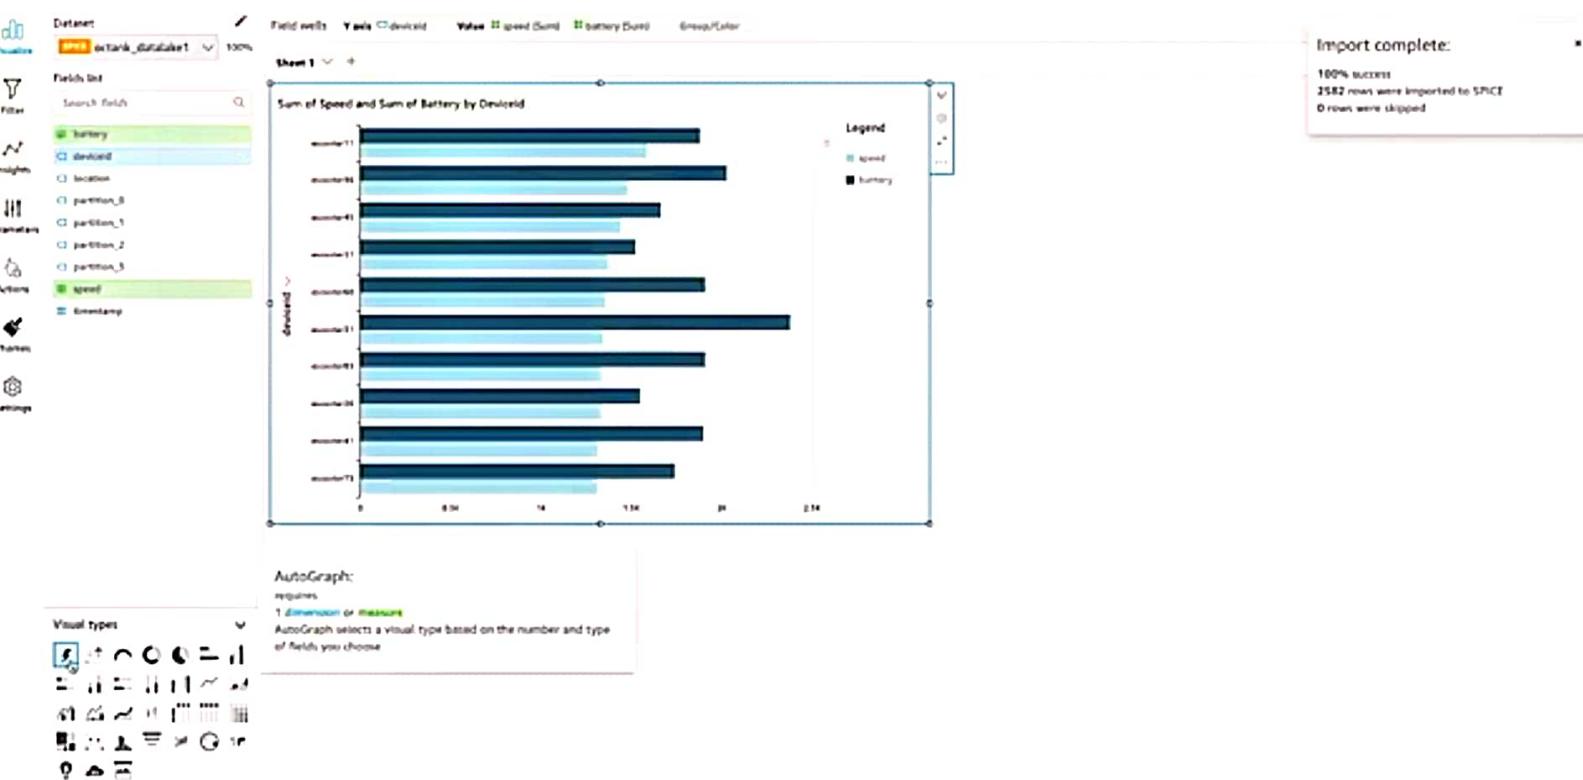
Choose 1 or more fields and let QuickSight choose the most appropriate chart



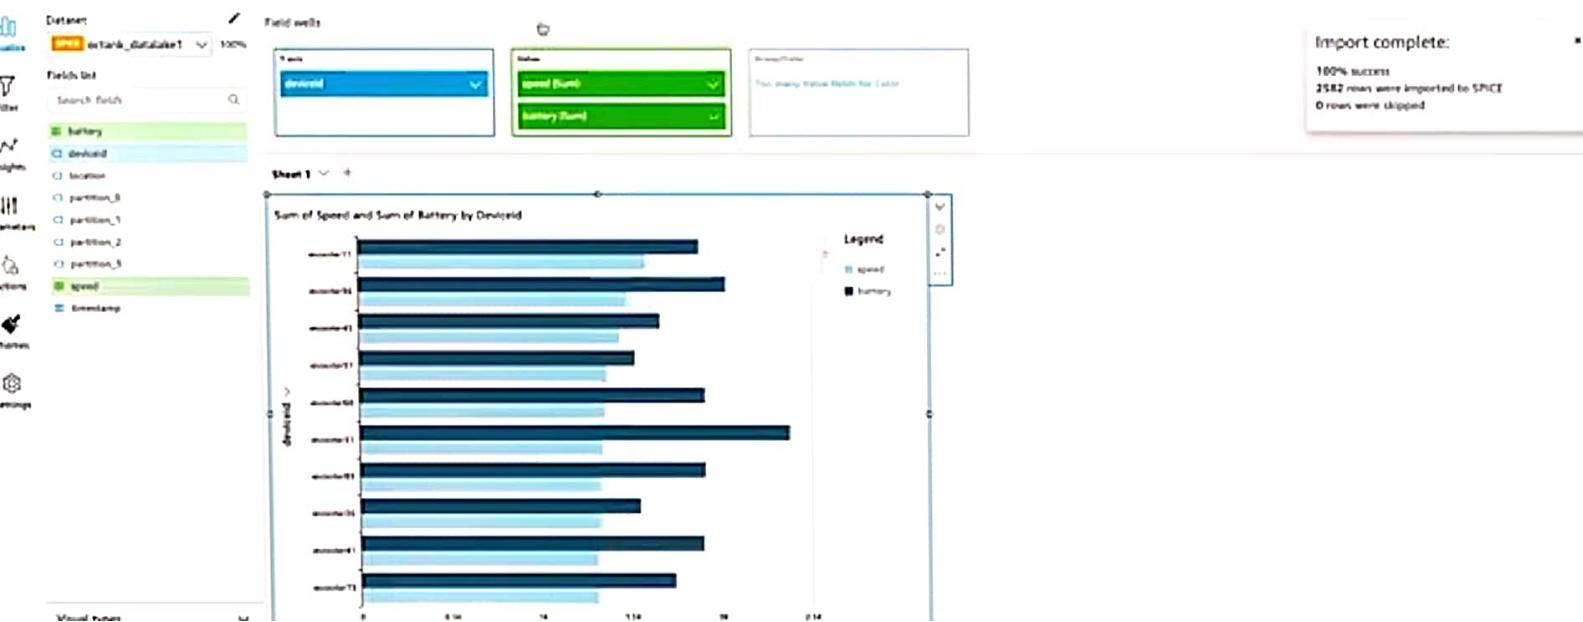


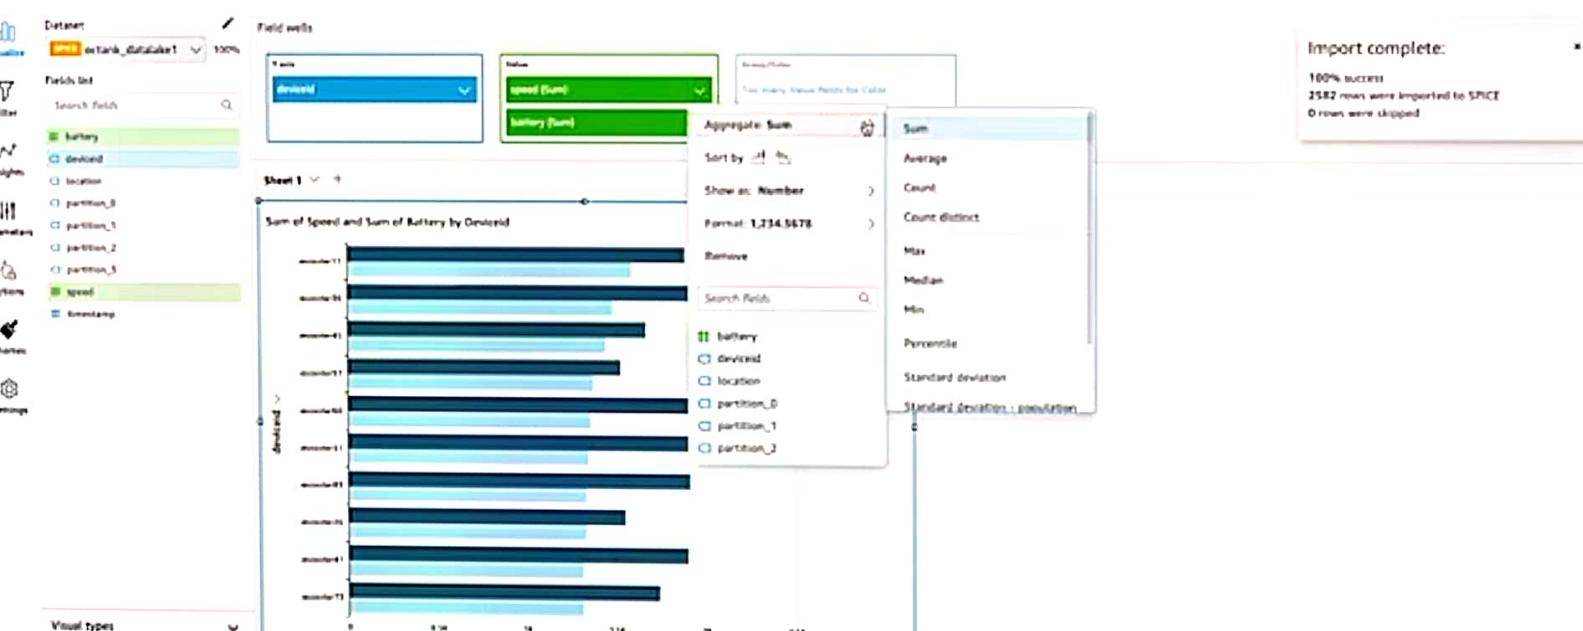


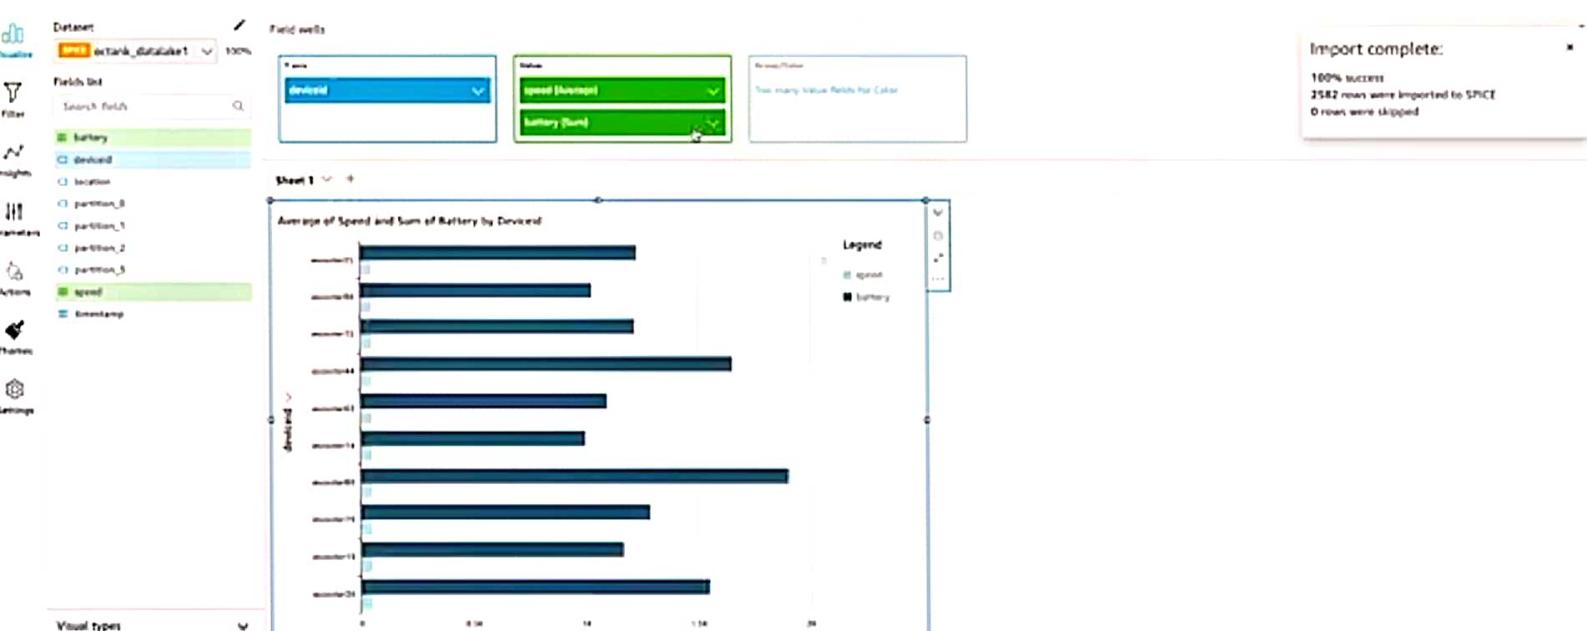


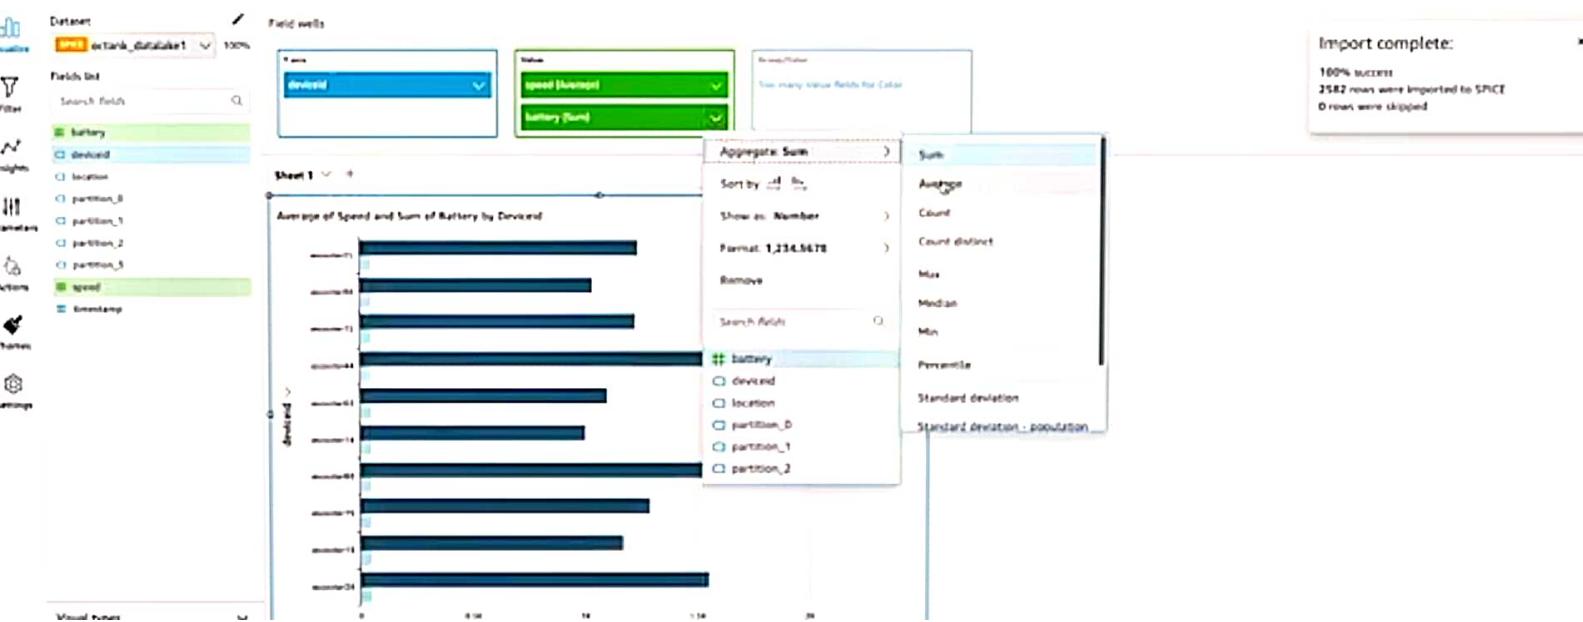








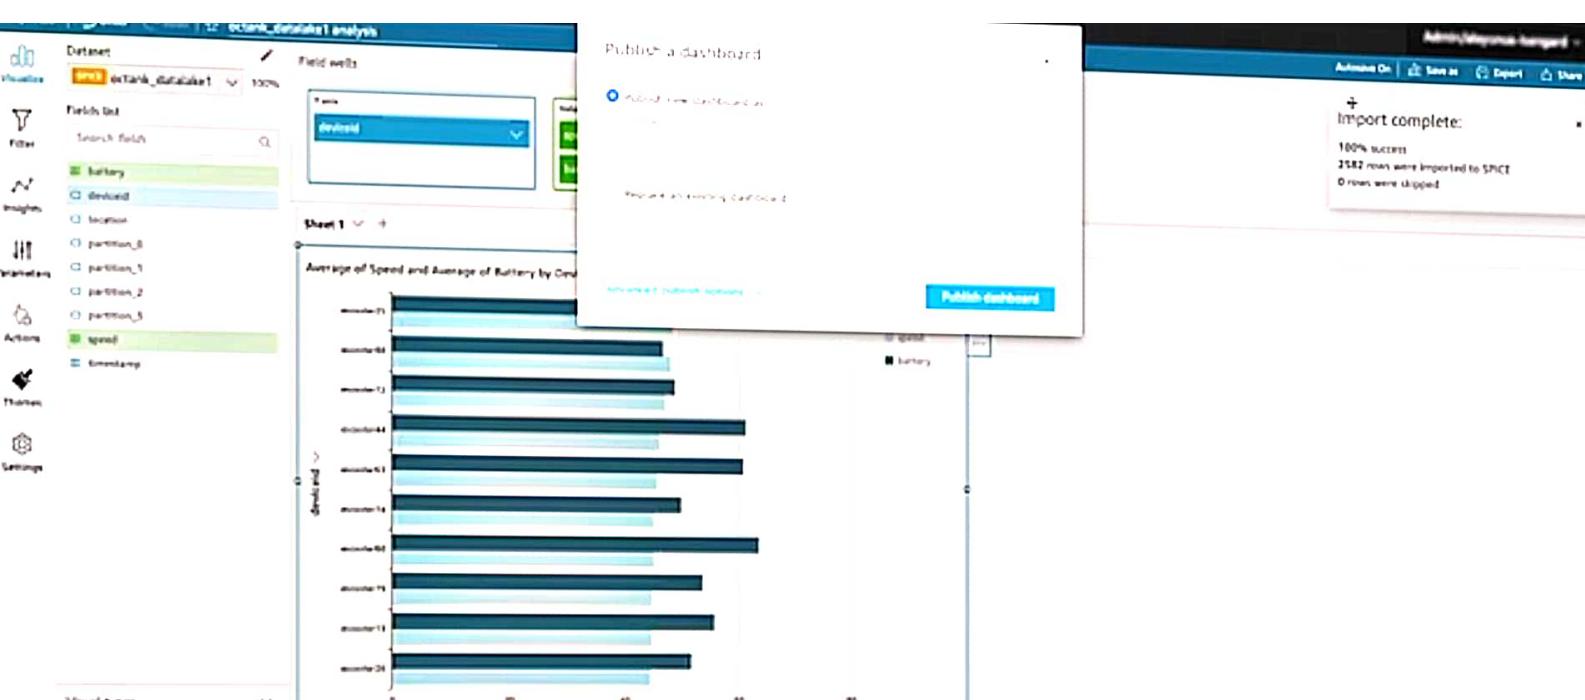








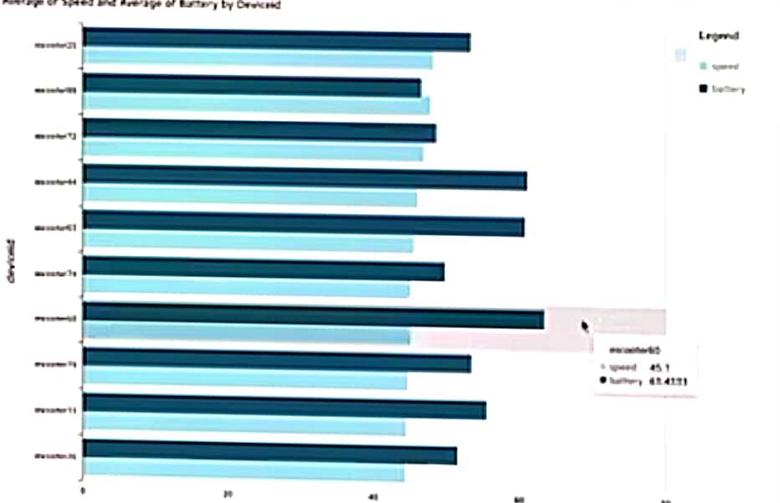












Pipeline to build

